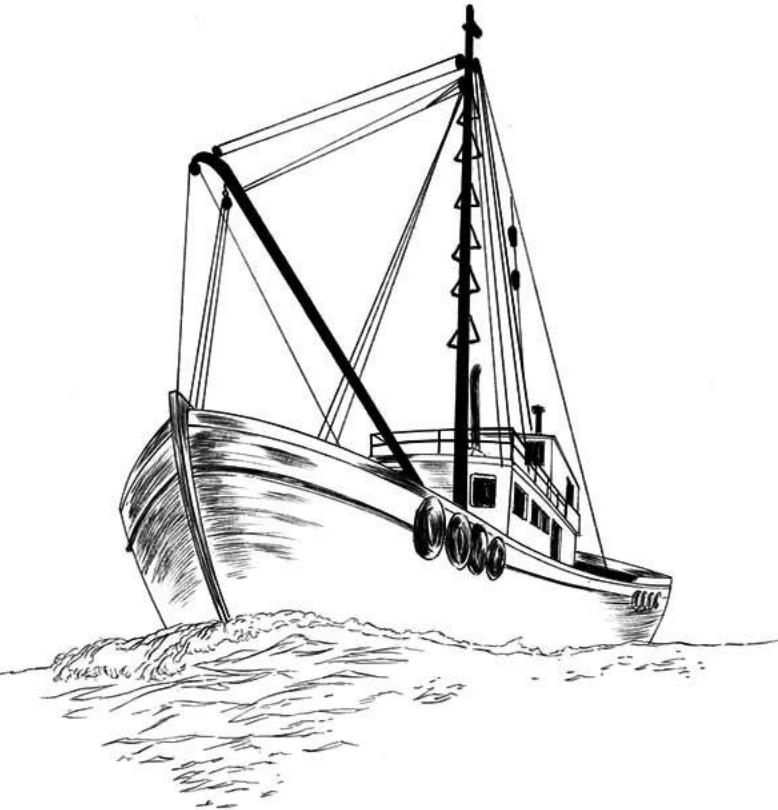




Before We Begin. What's the problem? Why Docker?



AGENDA

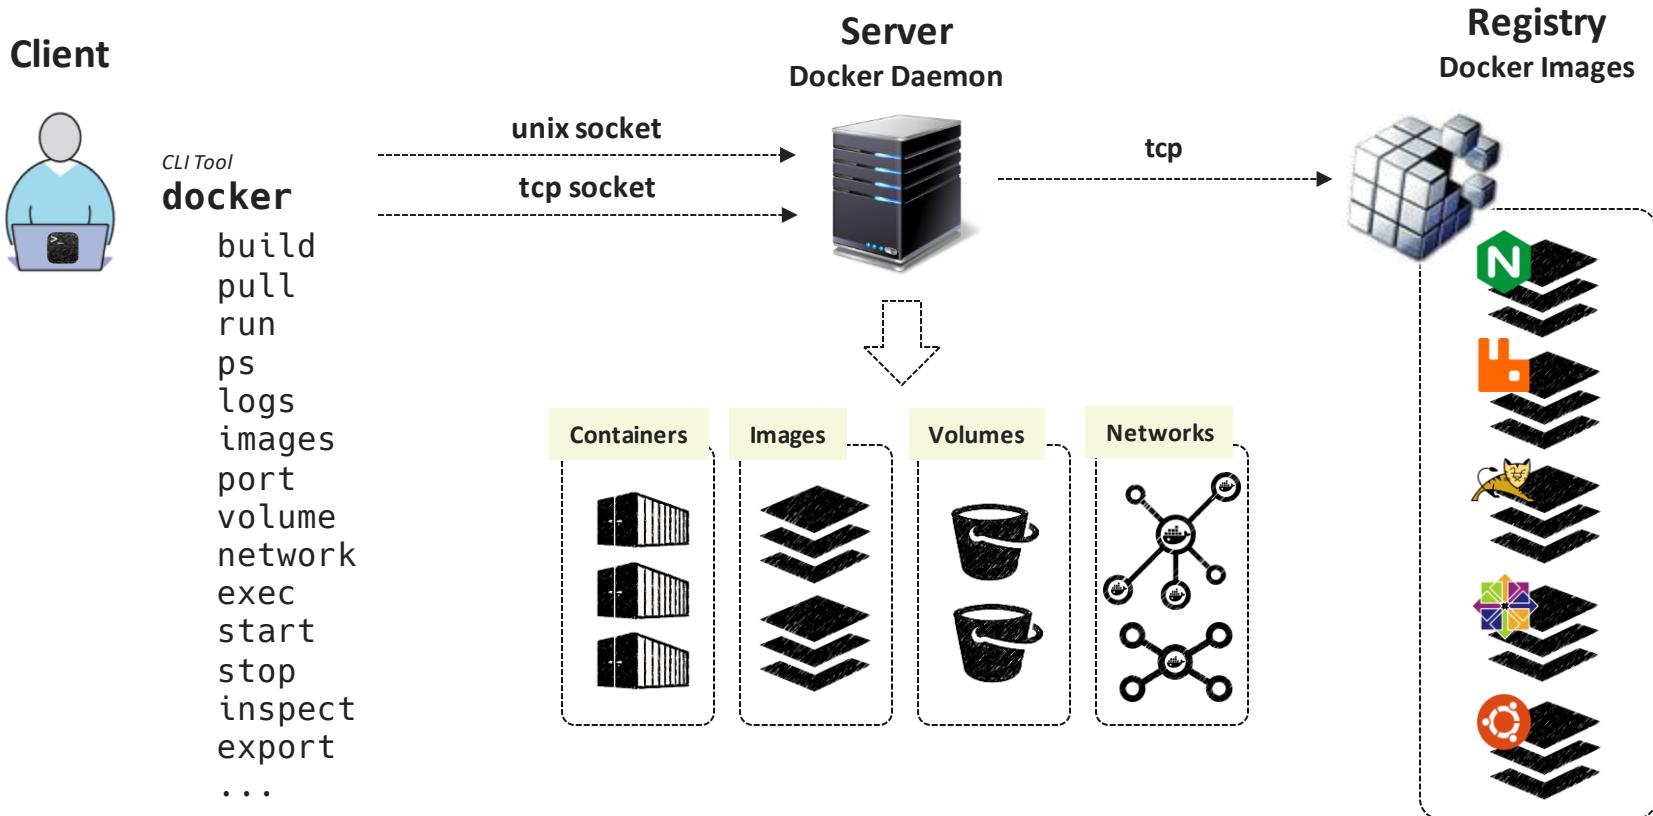
- Docker Architecture Overview
- Docker Installation and Configuration
- Creating Docker Images
- Running Containers
- Mounting Data from the Host



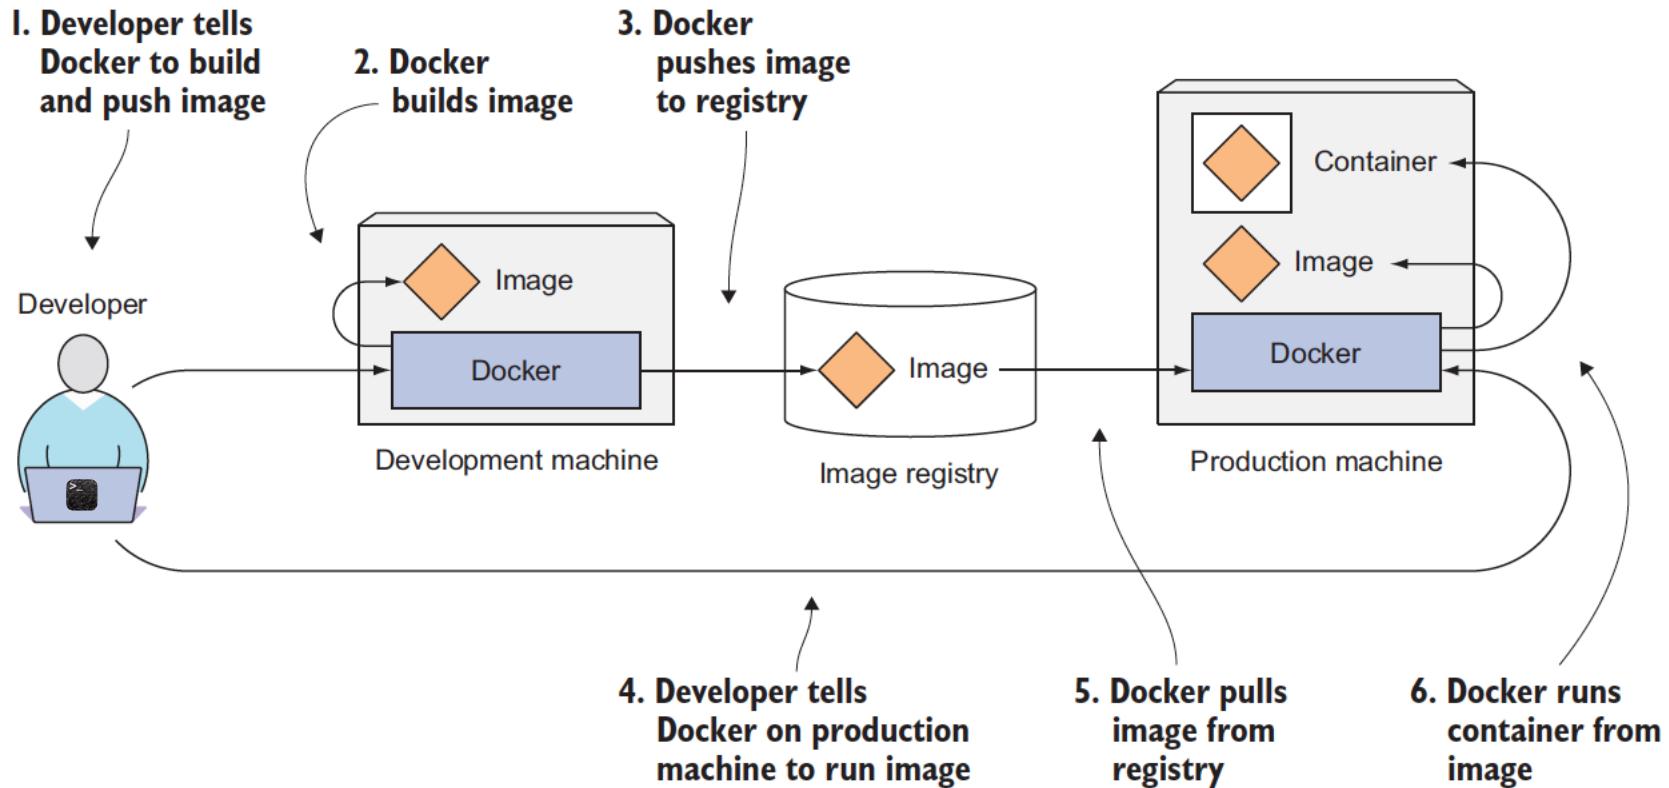
ARCHITECTURE OVERVIEW

- **Docker Architecture**
- **Docker Components**
- **Isolation Technologies**
- **What is a Container?**
- **What is an Image?**

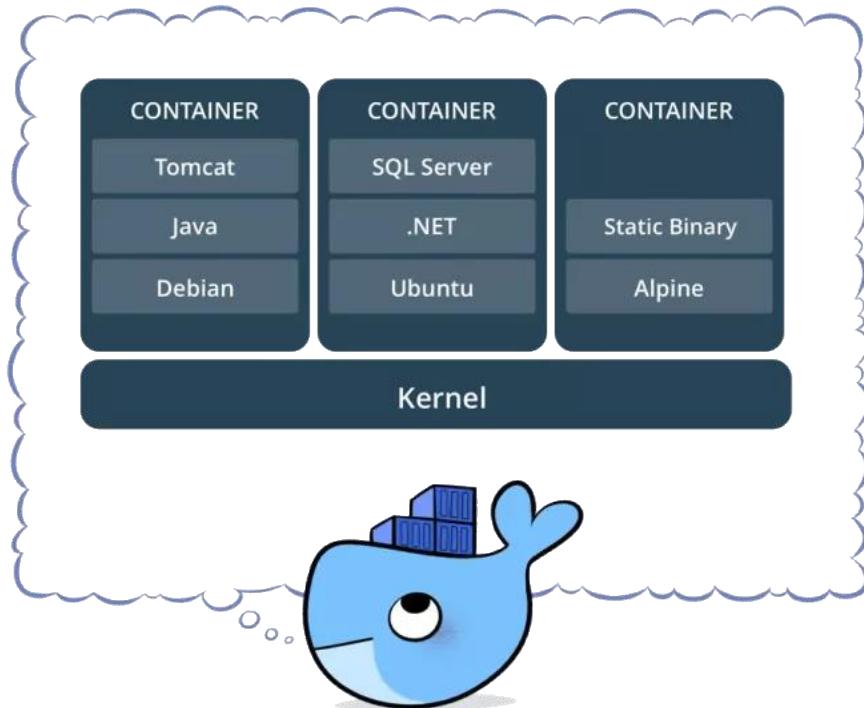
Docker Architecture



Docker Workflow



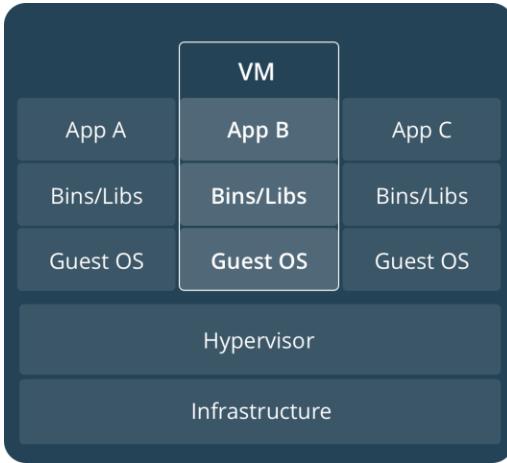
What is a Container?



A **container** is something quite similar to a virtual machine, which can be used to **contain and execute** all the software required to run a particular program or set of programs.

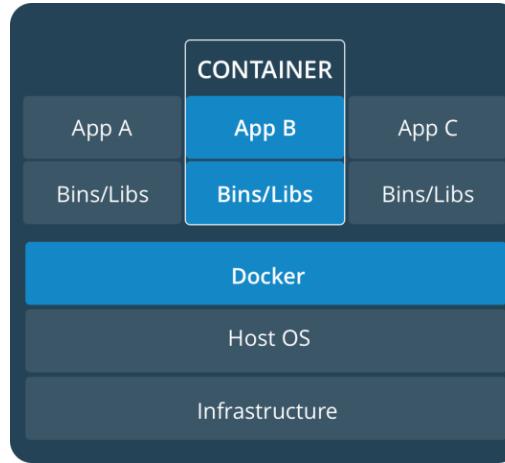
The container includes an operating system (typically some flavor of Linux) as base, plus any software installed on top of the OS that might be needed. This container can therefore be run as a self-contained virtual environment, which makes it a lot easier to reproduce the same analysis on any infrastructure that supports running the container, from your laptop to a cloud platform, without having to go through the pain of identifying and installing all the software dependencies involved. You can even have multiple containers running on the same machine, so you can easily switch between different environments if you need to run programs that have incompatible system requirements.

Virtual Machines vs Containers



VIRTUAL MACHINES

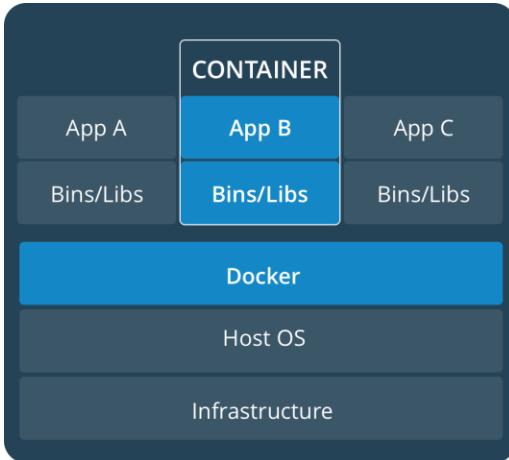
Virtual machines (VMs) are an abstraction (emulation) of physical hardware turning one server into many servers. The hypervisor allows multiple VMs to run on a single machine. Each VM includes a full copy of an operating system, one or more apps, necessary binaries and libraries - taking up tens of GBs. VMs can also be slow to boot.



CONTAINERS

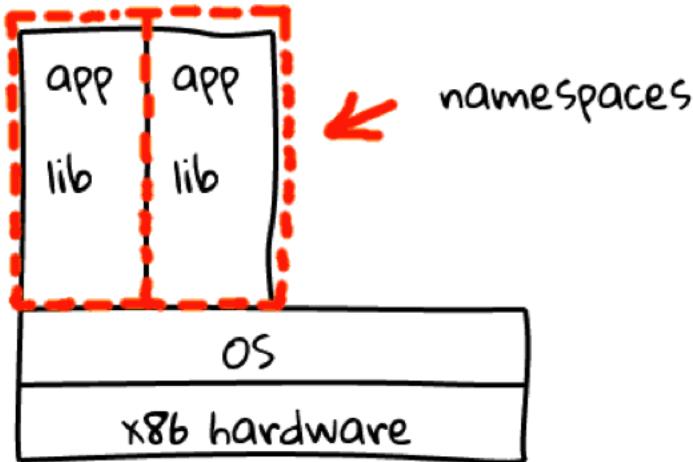
Containers are an abstraction at the app layer that packages code and dependencies together (worker process). Multiple containers can run on the same machine and share the OS kernel with other containers, each running as isolated processes in user space. Containers take up less space than VMs (container images are typically tens of MBs in size), and start almost instantly.

Isolation Technologies: Functional Requirements



Why	How
A process in the container shouldn't see other processes running in the system and should feel itself as the only process in the system	Namespaces
We should have facility to manage systems resources –(for example) to provide necessary amount of RAM, CPU shares and so on to the Container	Control Groups
A process in the container should see/use the only that part of the filesystems which is required by it	Chroot
A process in the container should have enough permissions to manage/use kernel	Process Capabilities
A process in the container should have access to ethernet device	Virtual eth
A few containers can expose the same ports, but it shouldn't cause any problem	Port Binding
A service should have facility to keep data, even if it goes down accidentally	Volumes
Service should be able to communicate with other services by IP/names	Docker Network

Isolation Technology



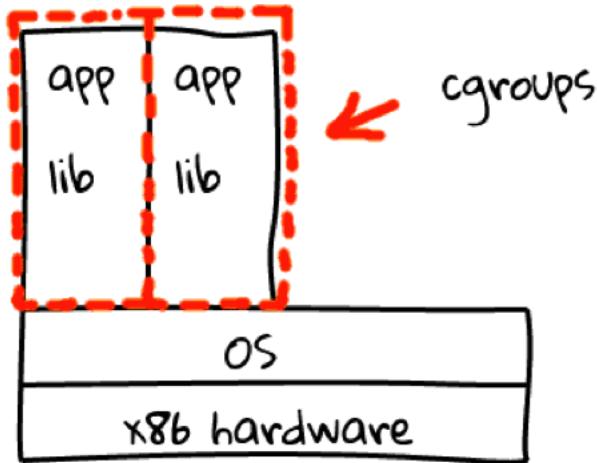
Namespaces

When you run a container, Docker creates a set of *namespaces* for that container. This provides a layer of isolation: each aspect of a container runs in its own namespace and does not have access outside of it.

There are 6 Linux namespaces that Docker Engine uses:

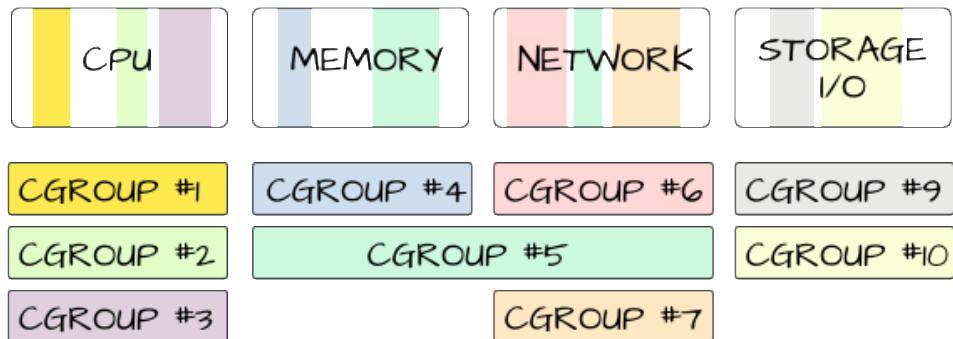
- ✓ The **mnt** namespace: Managing mount-points. *Linux 2.4.19*
- ✓ The **uts** namespace: Isolating kernel and version identifiers. (UTS: Unix Timesharing System). *Linux 2.6.19*
- ✓ The **ipc** namespace: Managing access to IPC resources (IPC: Inter Process Communication). *Linux 2.6.19*
- ✓ The **pid** namespace: Process isolation. *Linux 2.6.24*
- ✓ The **net** namespace: Managing network interfaces, *Linux 2.6.24*
- ✓ The **user** namespace. *Linux 3.8*

Isolation Technology

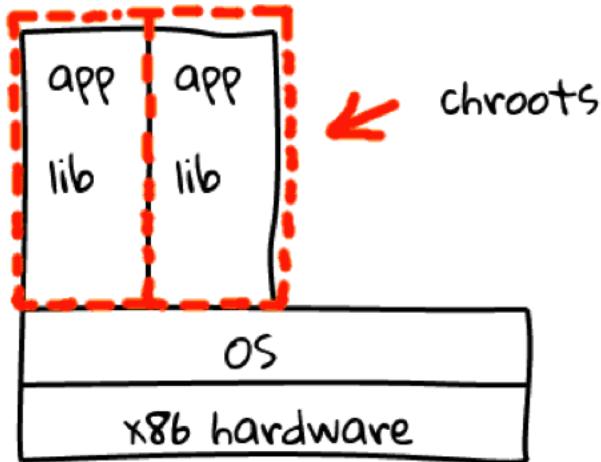


Control Groups

A key to running applications in isolation is to have them only use the resources you want. This ensures containers are good multi-tenant citizens on a host. Control groups allow Docker Engine to share available hardware resources to containers and, if required, set up limits and constraints. For example, limiting the memory available to a specific container.

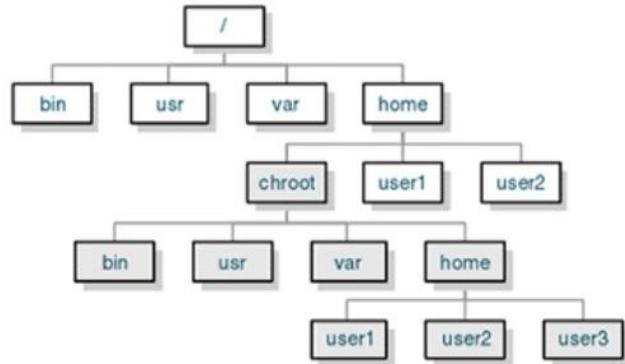


Isolation Technology



chroot

is simply isolation
on the filesystem

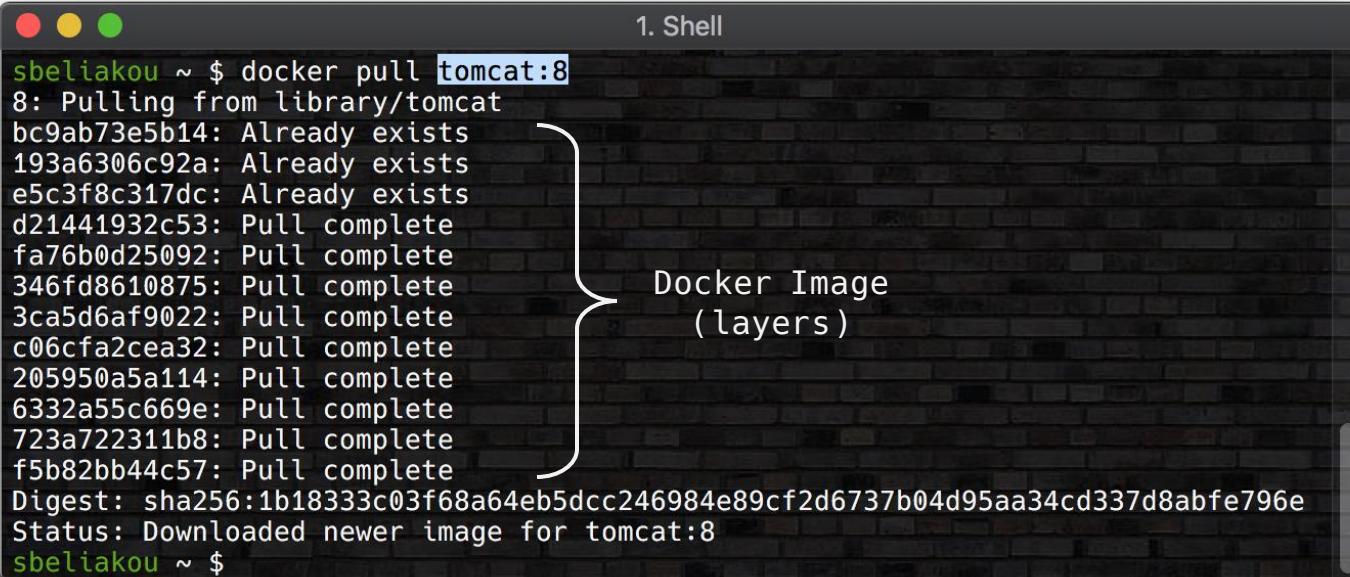


Union FS

operate by creating layers, making them very lightweight and fast. Docker Engine uses union file systems to provide the building blocks for containers. Docker Engine can make use of several union file system variants including: AUFS, btrfs, vfs, and DeviceMapper

Docker Image. What is this?

An image is an inert, immutable, file that's essentially a snapshot of a container. Images are created with the build command, and they'll produce a container when started with run. Images are stored in a Docker registry such as <https://hub.docker.com>. Because they can become quite large, images are designed to be composed of layers of other images, allowing a minimal amount of data to be sent when transferring images over the network



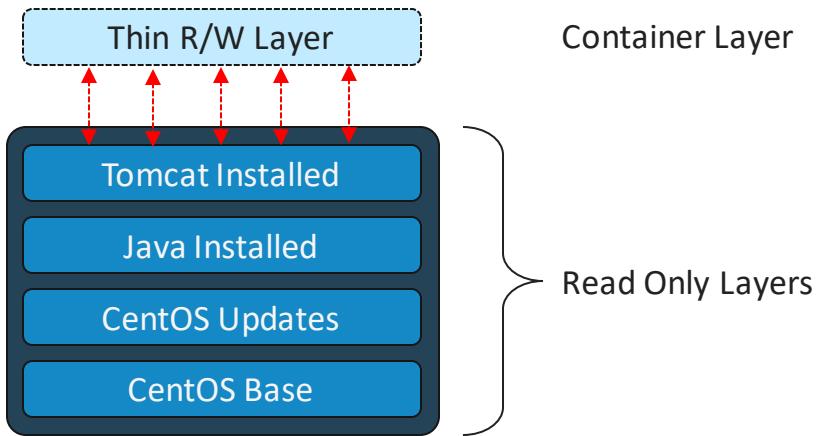
1. Shell

```
sbeliakou ~ $ docker pull tomcat:8
8: Pulling from library/tomcat
bc9ab73e5b14: Already exists
193a6306c92a: Already exists
e5c3f8c317dc: Already exists
d21441932c53: Pull complete
fa76b0d25092: Pull complete
346fd8610875: Pull complete
3ca5d6af9022: Pull complete
c06cfa2cea32: Pull complete
205950a5a114: Pull complete
6332a55c669e: Pull complete
723a722311b8: Pull complete
f5b82bb44c57: Pull complete
Digest: sha256:1b18333c03f68a64eb5dcc246984e89cf2d6737b04d95aa34cd337d8abfe796e
Status: Downloaded newer image for tomcat:8
sbeliakou ~ $
```

Docker Image
(layers)

Docker Storage (Graph) Drivers

UnionFS allows files and directories of separate file systems, known as branches, to be transparently overlaid, forming a single coherent file system. Contents of directories which have the same path within the merged branches will be seen together in a single merged directory, within the new, virtual filesystem.



The copy-on-write (CoW) strategy

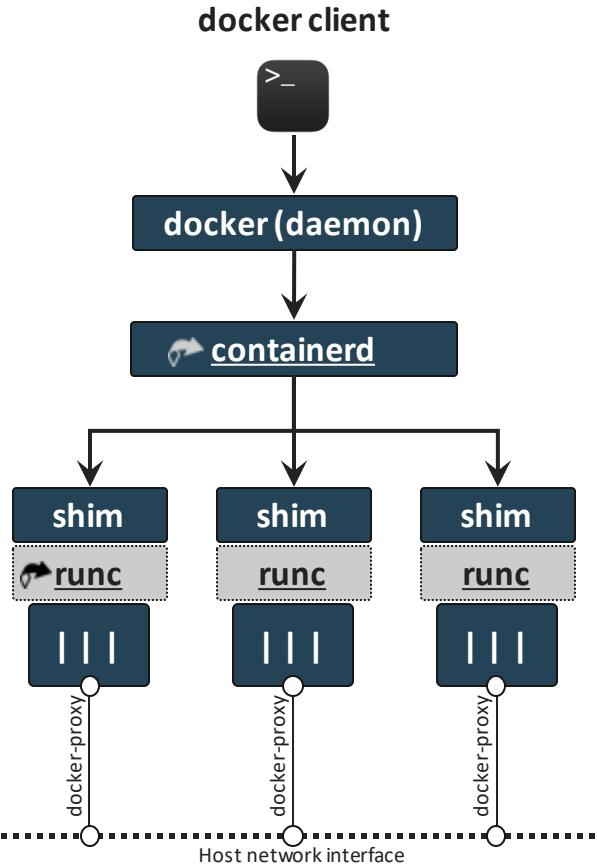
is a strategy of sharing and copying files for maximum efficiency. If a file or directory exists in a lower layer within the image, and another layer (including the writable layer) needs read access to it, it just uses the existing file. The first time another layer needs to modify the file (when building the image or running the container), the file is copied into that layer and modified. This minimizes I/O and the size of each of the subsequent layers.

Docker Graphdrivers

Given the “image graph” of layer content represents the relationships between various layers, the driver to handle these layers is called a “graphdriver.”

- [**VFS**](#) - does not use a Union FS or CoW, is very valuable for simple validation and testing of other parts of the Docker engine
- [**AUFS**](#) - enables shared memory pages for different containers loading the same shared libraries from the same layer (because they are the same inode on disk). Currently it's available on Debian and Ubuntu and not considered as mainstream.
- [**Overlay2**](#) - resolves the inode exhaustion problem as well as a few other bugs that were inherent to the old design of the original driver. It also enables shared memory between disparate containers using the same on-disk shared libraries.
- [**Device Mapper**](#) - It is quite unlike the union filesystems in that devicemapper works on block devices. There is no way to get default “out of the box” performance with devicemapper. Some of the features rely on specific versions of libdevmapper and it requires above-average skill to validate all these settings on a system.
- [**Btrfs**](#) - a disk formatted as a btrfs filesystem is required as the graphdriver root (by default, /var/lib/docker). It's not supported by RedHat.

Docker Infrastructure



dockerd:

The Docker daemon itself. The highest level component in your list and also the only 'Docker' product listed. Provides all the nice UX features of Docker.

docker-containerd:

It's a [daemon](#), listening on a Unix socket, exposes gRPC endpoints. Handles all the low-level container management tasks, storage, image distribution, network attachment, etc...

docker-containerd-ctr:

A lightweight CLI to directly communicate with containerd. Think of it as how 'docker' is to 'dockerd'.

docker-containerd-shim:

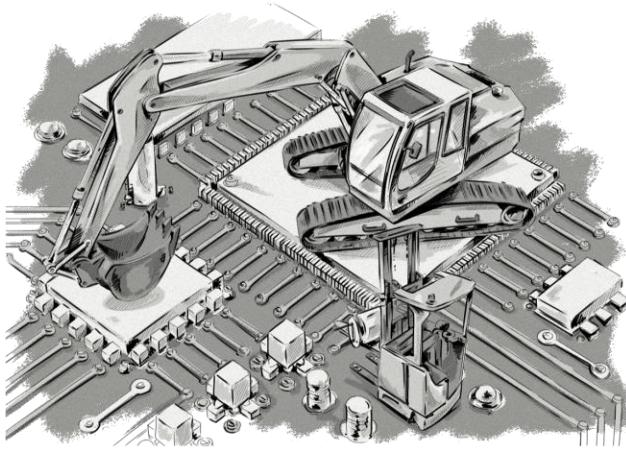
After [runC](#) actually runs the container, it exits (allowing us to not have any long-running processes responsible for our containers). The shim is the component which sits between containerd and runc to facilitate this.

docker-runc:

A lightweight binary for actually running containers. Deals with the low-level interfacing with Linux capabilities like cgroups, namespaces, etc...

docker-proxy:

A tool responsible for proxying container's ports to Host's interface



Docker installation and Configuration

- **Installing Docker on CentOS**
- **How Docker Works**
- **Configuring User Privileges to Talk to Dockerd**
- **Configuring Docker Daemon**

Learning Docker

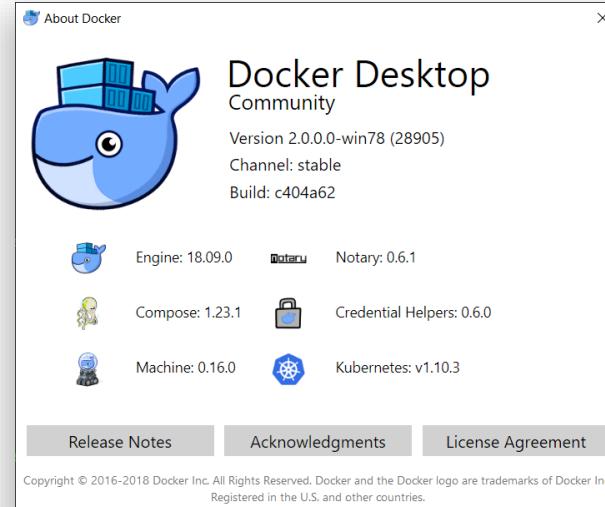
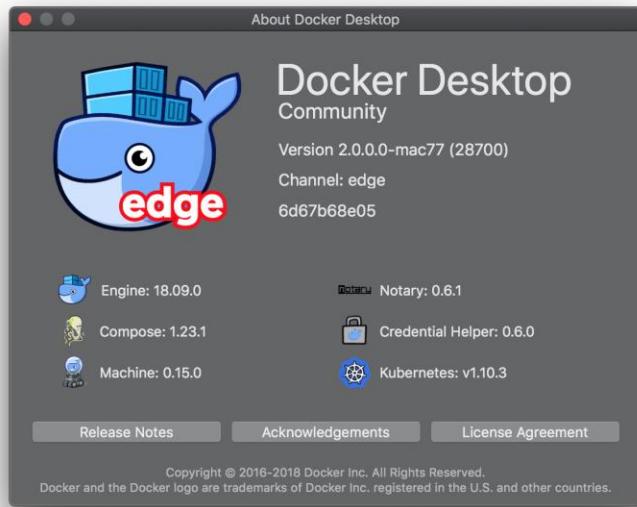


- ➡ [Install Docker CE on CentOS](#)
- ➡ [Install Docker CE on Ubuntu](#)

- ➡ [Play with Docker Lab](#)
- ➡ [Play with Docker Classroom](#)

- ➡ [Docker Scenarios on Katacoda](#)

Docker on MacOS X and Windows

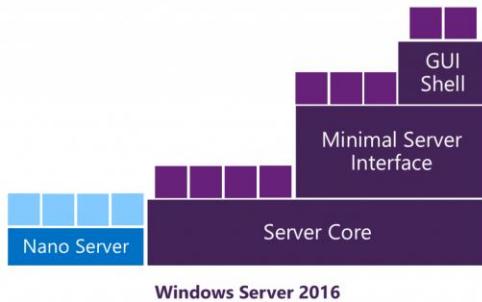


 [Install Docker for Mac](#)

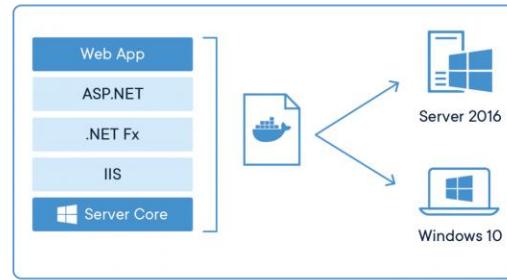
 [Install Docker for Windows](#)
 [docs.microsoft.com/...](https://docs.microsoft.com/)

Docker on Windows

Microsoft includes two different types of containers:



“Hyper-V Container”,
based on the *Windows Nano Server* image



“Windows Server Container”,
based on *Windows Server Core* image

Check out more details here:

- [windows-containers](#)
- [quick-start-windows-10](#)

And here:

- [working-windows-containers-docker-basics](#)
- [working-windows-containers-docker-running](#)
- [working-windows-containers-docker-stride](#)

Local Engineering Environment

[Get started with VAGRANT](#)



VAGRANT



Microsoft
Hyper-V

A screenshot of a code editor window titled "Vagrantfile". The file contains Ruby code for configuring a Vagrant environment:

```
1 Vagrant.configure("2") do |config|
2   config.vm.box = "sbeliakou/centos"
3   config.vm.network :private_network, ip: "192.168.56.15"
4
5   config.vm.provision "shell", inline: <<-SHELL
6     yum install -y yum-utils jq net-tools
7     yum-config-manager --add-repo \
8       https://download.docker.com/linux/centos/docker-ce.repo
9     yum-config-manager --enable docker-ce-edge
10
11    yum install -y docker-ce
12    systemctl enable docker
13    systemctl start docker
14    usermod -aG docker vagrant
15
16 SHELL
17 end
```

The code provisions a CentOS VM with Docker CE, installs necessary packages, and adds the vagrant user to the docker group. A link to the GitHub repository is shown at the bottom right: <https://github.com/sbeliakou/docker-training>.

A screenshot of a terminal window titled "1. Shell". It shows examples of Vagrant commands:

```
sbeliakou ~/project_folder $ vagrant up      # create a VM from Vagrantfile
sbeliakou ~/project_folder $ vagrant ssh        # ssh in to VM
sbeliakou ~/project_folder $ vagrant halt       # stop VM
sbeliakou ~/project_folder $ vagrant destroy    # destroy VM
```

Installing Docker Service

The docker daemon binds to a Unix socket /var/run/docker.sock which is owned by root:docker.

```
[vagrant@localhost ~]$ docker info  
Got permission denied while trying to connect to the Docker daemon socket at  
unix:///var/run/docker.sock: Get http://%2Fvar%2Frun%2Fdocker.sock/v1.38/containers/json:  
dial unix /var/run/docker.sock: connect: permission denied
```

Non-root user just needs to be added to the docker group.

```
[vagrant@localhost ~]$ sudo usermod -aG docker vagrant  
# log out / log in  
  
[vagrant@localhost ~]$ id  
uid=1000(vagrant) gid=1000(vagrant) groups=1000(vagrant),993(docker)
```

Let's check docker info from vagrant user

```
[vagrant@localhost ~]$ docker info  
Containers: 0  
Running: 0  
Paused: 0  
Stopped: 0  
...  
...
```

How It Works

Either by using the docker binary or via the API, the Docker client tells the Docker daemon to run a container.

```
[vagrant@localhost ~]$ docker run -i -t ubuntu bash
```

The Docker Engine client is launched using the *docker* tool with the *run* option running a new container.

The bare minimum the *docker client* needs to tell the Docker daemon to run the container is:

- ✓ What Docker image to build/create the container from, for example, *ubuntu*
- ✓ The command you want to run inside the container when it is launched, for example, */bin/bash*

So what happens under the hood when we run this command?

In order, Docker Engine does the following:

- ✓ Pulls the *ubuntu* image.
- ✓ Creates a new container.
- ✓ Allocates a filesystem and mounts a read-write *layer*.
- ✓ Allocates a network / bridge interface.
- ✓ Sets up an IP address.
- ✓ Executes a process that you specify.
- ✓ Captures and provides application output.

You now have a running container! Now you can manage your container, interact with your application and then, when finished, stop and remove your container.

Try Something Simple

```
[vagrant@localhost vagrant]$ docker run hello-world  
Unable to find image 'hello-world:latest' locally  
latest: Pulling from library/hello-world  
9db2ca6ccae0: Pull complete  
Digest: sha256:4b8ff392a12ed9ea17784bd3c9a8b1fa3299cac44aca35a85c90c5e3c7afacdc  
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker!

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker ID:

```
https://hub.docker.com/
```

For more examples and ideas, visit:

```
https://docs.docker.com/engine/userguide/
```

Docker Daemon, Hub/Registry Commands

Display Docker Version and Info

- `docker --version`
- `docker version`
- `docker version --format '{{.Server.Version}}'`
- `docker version --format '{{json .}}'`
- `docker info`
- `docker info --format '{{json .}}'`

Working with Docker Hub/Registry:

- [docker login](#) - to login to a registry.
- [docker logout](#) - to logout from a registry.
- [docker search](#) - searches registry for image.
- [docker pull](#) - pulls an image from registry to local machine.
- [docker push](#) - pushes an image to the registry from local machine.

Inspecting Docker Configuration

```
[vagrant@localhost ~]$ docker info
Containers: 0
Running: 0
Paused: 0
Stopped: 0
Images: 0
Server Version: 18.06.0-ce
Storage Driver: overlay2
    Backing Filesystem: xfs
Logging Driver: json-file
Runtimes: runc
Default Runtime: runc
Kernel Version: 3.10.0-862.9.1.el7.x86_64
Operating System: CentOS Linux 7 (Core)
Docker Root Dir: /var/lib/docker
Registry: https://index.docker.io/v1/
...
```

```
[vagrant@localhost]$ ps -eo %a | grep [d]ocker
/usr/bin/dockerd
docker-containerd --config /var/run/docker/containerd/containerd.toml
```



Docker Root Dir

```
# tree -ACFL 1 /var/lib/docker
/var/lib/docker
├── builder/
├── buildkit/
├── containerd/
├── containers/
├── image/
├── network/
├── overlay2/
├── plugins/
├── runtimes/
├── swarm/
├── tmp/
├── trust/
└── volumes/
```

Configuring Docker Daemon

```
# docker info | egrep "(Cgroup|Storage) Driver"
Storage Driver: overlay2
Cgroup Driver: cgroupfs
```

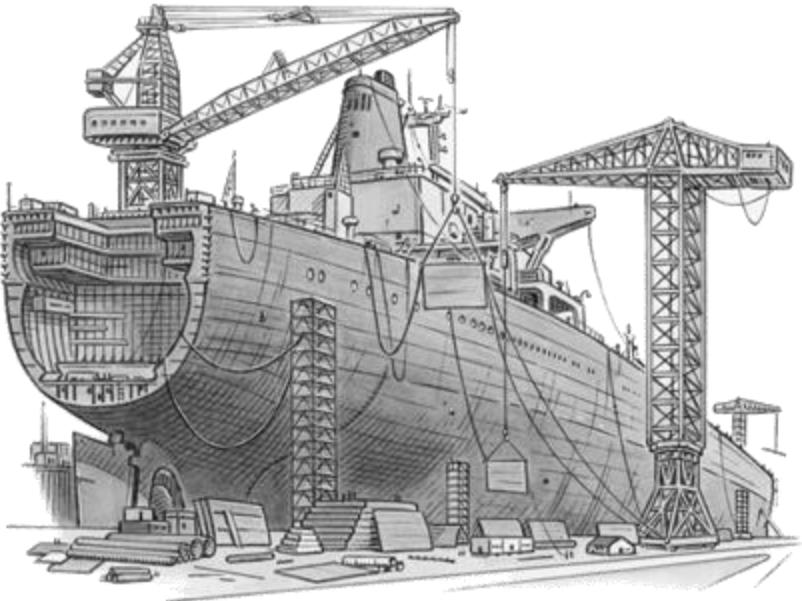
/etc/docker/daemon.json

```
{
  "exec-opts": [
    "native.cgroupdriver=systemd"
  ],
  "storage-driver": "devicemapper"
}
```

```
# systemctl daemon-reload
# systemctl restart docker.service
# docker info | egrep "(Cgroup|Storage) Driver"
Storage Driver: devicemapper
Cgroup Driver: systemd
```



[More Configuration Options](#)



Creating Docker Images

- Build own Image with Dockerfile
- Tagging Images
- Entrypoint and CMD
- Build Arguments
- Multistage Build
- ONEBUILD Instructions
- Health Checks

Working with Images

Lifecycle:

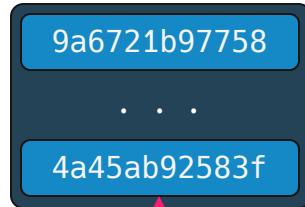
- [`docker images`](#) - shows all images.
- [`docker import`](#) - creates an image from a tarball.
- [`docker build`](#) - creates image from Dockerfile.
- [`docker commit`](#) - creates image from a container, pausing it temporarily if it is running.
- [`docker rmi`](#) - removes an image.
- [`docker load`](#) - loads an image from a tar archive as STDIN, including images and tags.
- [`docker save`](#) - saves an image to a tar archive stream to STDOUT with all parent layers, tags & versions.

Info:

- [`docker history`](#) - shows history of image.
- [`docker tag`](#) - tags an image to a name (local or registry).

What is an Image?

trainings:centos-node



ansible:2.6.2

639b4cd19004

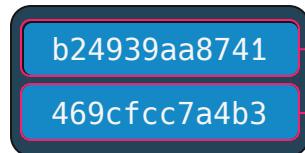
ansible:2.6.1

a3c99315580f

22a01822f3c0

f52300eb0803

sbeliakou/centos:latest



```
@bb6140f107a7:~  
sbeliakou ~ $ docker pull sbeliakou/centos:latest  
latest: Pulling from sbeliakou/centos  
469cfcc7a4b3: Pull complete  
b24939aa8741: Pull complete  
Digest: sha256:46bfd372002e25fb3505b0119853bdfbfee26395cdb97af45ead2aa83c685394  
Status: Downloaded newer image for sbeliakou/centos:latest  
sbeliakou ~ $ docker pull sbeliakou/ansible:2.6.1  
2.6.1: Pulling from sbeliakou/ansible  
469cfcc7a4b3: Already exists  
f52300eb0803: Pull complete  
22a01822f3c0: Pull complete  
a3c99315580f: Pull complete  
Digest: sha256:ce24829798b626c8f748dc225cbaf805189cd3945f9a37f89a53f3da393baf2  
Status: Downloaded newer image for sbeliakou/ansible:2.6.1  
sbeliakou ~ $ docker pull sbeliakou/ansible:2.6.2  
2.6.2: Pulling from sbeliakou/ansible  
469cfcc7a4b3: Already exists  
f52300eb0803: Already exists  
22a01822f3c0: Already exists  
a3c99315580f: Already exists  
639b4cd19004: Pull complete  
Digest: sha256:46904da2ce545c753f14b13e73d8d2097fe431c9ca013ac0e08e4c0e6fd8d247  
Status: Downloaded newer image for sbeliakou/ansible:2.6.2  
sbeliakou ~ $ docker pull sbeliakou/trainings:centos-node  
centos-node: Pulling from sbeliakou/trainings  
469cfcc7a4b3: Already exists  
b24939aa8741: Already exists  
4a45ab92583f: Pull complete  
b607a6c71319: Pull complete  
406aff0ed91f: Pull complete  
7e5f845fcf24: Pull complete  
40e8e164788e: Pull complete  
e96b1ac125c0: Pull complete  
b9debd79786: Pull complete  
20439a1df421: Pull complete  
5565437f488a: Pull complete  
1fb62f64a7ab: Pull complete  
a194ef741153: Pull complete  
7d5d09b24dcdb: Pull complete  
cd9bc492f001: Pull complete  
cafa18c0e0f3: Pull complete  
9a6721b97758: Pull complete  
Digest: sha256:4c7c2b3ce1c22ddd99b377a1ac12413389fcf4b703ce9f24694ea3d5f5a4c205  
Status: Downloaded newer image for sbeliakou/trainings:centos-node  
sbeliakou ~ $
```

Writing Own Dockerfile

→ <http://flask.pocoo.org/>

→ <http://flask.pocoo.org/docs/1.0/quickstart/#quickstart>



```
from flask import Flask
app = Flask(__name__)

@app.route("/")
def hello():
    return "Hello World!\n"
```

>-

```
$ docker build -t myflask .
Successfully built 23bbddc9f137
Successfully tagged myflask:latest
$ docker run -d -P myflask
6644fd128502...
$ docker port 6644fd128502
5000/tcp -> 0.0.0.0:32779
$ curl localhost:32779
```

→ [dockerfile_best-practices](#)

→ [sbeliakou/flask-hello](#)



Dockerfile

```
FROM python:2.7

RUN pip install Flask
COPY . / 

ENV FLASK_APP=hello.py
EXPOSE 5000
CMD flask run --host=0.0.0.0
```

Each instruction creates one layer:

- **FROM** creates a layer from the python:2.7 Docker image.
- **RUN** performs necessary commands
- **COPY** adds files from your Docker client's current directory
- **ENV** sets environment variable
- **EXPOSE** defines which ports will be exposed by container
- **CMD** specifies what command to run within the container

Dockerfile Instructions

- .dockerignore
- FROM - Sets the Base Image for subsequent instructions.
- RUN - execute any commands in a new layer on top of the current image and commit the results.
- CMD - provide defaults for an executing container.
- EXPOSE - informs Docker that the container listens on the specified network ports at runtime.
NOTE: does not actually make ports accessible.
- ENV - sets environment variable.
- ADD - copies new files, directories or remote file to container. Invalidates caches.
- COPY - copies new files or directories to container.
- ENTRYPOINT - configures a container that will run as an executable.
- VOLUME - creates a mount point for externally mounted volumes or other containers.
- USER - sets the user name for following RUN / CMD / ENTRYPOINT commands.
- WORKDIR - sets the working directory.
- ARG - defines a build-time variable.
- ONBUILD - adds a trigger instruction when the image is used as the base for another build.
- STOPSIGNAL - sets the system call signal that will be sent to the container to exit.
- LABEL - apply key/value metadata to your images, containers, or daemons.

Useful Links and Examples

Docker CLI Reference: [commandline/cli/](#)

Dockerfile Reference: [#/dockerfile-reference](#)

Many Samples: [docs.docker.com/samples/](#)

Examples: [#dockerfile-examples](#)

Jenkins: [Dockerfile-alpine](#)

Tomcat: [jre8-alpine/Dockerfile](#)

OpenJDK: [jdk/alpine/Dockerfile](#)

AmazonLinux: [2018.03/Dockerfile](#)

Another example is here: [#define-a-container-with-dockerfile](#)

Best practices for writing Dockerfiles: [docs.docker.com/dockerfile_best-practices/](#)

Docker Lab Samples: <https://github.com/docker/labs>

Dockerfile Practice Questions: [docker-certification-practice-questions-dockerfile](#)

Many Docker based solutions: [veggiemonk/awesome-docker](#)

Dockerfile



```
FROM centos
LABEL maintainer="Siarhei Beliakou"
RUN yum install -y httpd web-assets-httpd && \
    yum clean all
RUN echo "my httpd container" > /var/www/html/index.html
# ADD/COPY index.html /var/www/html/
EXPOSE 80
ENTRYPOINT ["httpd"]
CMD ["-DFOREGROUND"]
```



[Dockerfile](#)

```
$ docker build [-t image_tag] [-f ./path/to/dockerfile] .
...
Status: Downloaded newer image for centos:latest
--> 49f7960eb7e4
...
Step 4/6 : RUN echo "my httpd container" > /var/www/html/index.html
--> Running in 37cc17740d0b
...
Successfully built 50a986f614d5
```



Tagging Images

```
$ docker build -t myhttpd .
```

```
...
```

```
Successfully built 50a986f614d5
```

```
Successfully tagged myhttpd:latest
```

```
$ docker tag << id | tag >> << new_tag >>
```

```
$ docker tag 50a986f614d5 sbeliakou/myhttpd:1.0
```

```
$ docker tag myhttpd sbeliakou/myhttpd:1.0
```

```
$ docker tag myhttpd:latest sbeliakou/myhttpd:latest
```

```
$ docker push sbeliakou/myhttpd:latest
```

ENTRYPOINT and CMD



```
...  
ENTRYPOINT ["httpd"]  
CMD ["-DFOREGROUND"]
```

- An **ENTRYPOINT** allows you to configure a container that will run as an executable.
- The main purpose of a **CMD** is to provide defaults for an executing container
- If you would like your container to run the same executable every time, then you should consider using **ENTRYPOINT** in combination with **CMD**
- **ENTRYPOINT/CMD** has 2 forms: exec and shell
 - exec form:** ["echo", "hello", "world"] ← preferred form
 - shell form:** echo hello world ← supports ENV Vars resolving

👉 <https://docs.docker.com/engine/reference/builder/#entrypoint>

👉 <https://docs.docker.com/engine/reference/builder/#cmd>

👉 <http://www.johnzaccone.io/entrypoint-vs-cmd-back-to-basics/>

ENTRYPOINT and CMD



```
FROM busybox  
ENTRYPOINT ["ping"]  
CMD ["-c1", "epam.com"]
```

```
$ docker build -t ping:1.0 .
```

```
$ docker run ping:1.0
```

will send 1 ICMP package to epam.com (default)

```
$ docker run ping:1.0 --help
```

will print help output

```
$ docker run ping:1.0 google.com
```

will ping "google.com" unless interrupted with ^C

```
$ docker run ping:1.0 -c1 google.com
```

will ping "google.com" for 1 time

Dockerfile Example: Custom Packer Image



```
FROM centos:7
ARG PACKER_VERSION=1.2.3
RUN yum install -y \
    epel-release \
    yum-plugin-ovl \
    wget unzip \
    rsync \
    openssh openssh-clients && \
    yum install -y python-pip && \
    yum clean all

RUN wget -q https://releases.hashicorp.com/...${PACKER_VERSION}_linux_amd64.zip && \
    unzip -q packer_${PACKER_VERSION}_linux_amd64.zip -d /bin/ && \
    rm -f packer_${PACKER_VERSION}_linux_amd64.zip
RUN pip install -U ansible ansible-modules-hashivault

RUN useradd packer

USER packer
ENV USER packer

ENTRYPOINT ["/bin/packer"]
CMD ["--version"]
```

A Few More Examples



```
FROM openjdk:8-jre  
  
ADD customer-contact-service.jar /  
EXPOSE 4040  
  
ENTRYPOINT ["java", "-jar", "/customer-contact-service.jar"]
```



```
FROM openjdk:8  
  
ENV PORT=8080  
EXPOSE 8080  
COPY wiremock /wiremock  
  
ENTRYPOINT ["/wiremock/bin/startServer.sh"]
```

Docker Base Images

- **scratch** – this is the ultimate base image and it has 0 files and 0 size.
- **busybox** – a minimal Unix weighing in at 2.5 MB and around 10000 files.
- **debian:jessie** – the latest Debian is 122 MB and around 18000 files.
- **alpine:latest** – Alpine Linux, only 8 MB in size and has access to a package repository



```
FROM scratch
```



[More details](#)

```
ADD centos.tar.gz /
```

```
RUN yum install -y epel-release && \
    yum update -y && \
    yum clean all
```

```
LABEL architecture="amd64" \
      OS="CentOS" \
      License=GPLv2 \
      maintainer="Siarhei Beliakou (sbeliakou@gmail.com)"
```

Default command

```
CMD ["/bin/bash"]
```

Building With Arguments



```
ARG BASE_IMAGE  
FROM ${BASE_IMAGE}  
...
```

```
# docker build --build-arg BASE_IMAGE=ubuntu:16.04 .
```

```
# docker build --build-arg BASE_IMAGE=ubuntu:18.04 .
```



```
ARG BUILD_NUMBER  
ARG JOB_NAME  
LABEL build_number="${BUILD_NUMBER}"  
LABEL job_name="${JOB_NAME}"
```

```
# docker build \  
--build-arg BUILD_NUMBER=${BUILD_NUMBER} \  
--build-arg JOB_NAME=${JOB_NAME} \  
.
```

Using ARG Directive

```
1  FROM openjdk:8-jdk
2
3  RUN apt-get update && apt-get install -y git curl && rm -rf /var/lib/apt/lists/*
4
5  ARG user=jenkins
6  ARG group=jenkins
7  ARG uid=1000
8  ARG gid=1000
9  ARG http_port=8080
10 ARG agent_port=50000
11 ARG JENKINS_HOME=/var/jenkins_home
...
19  RUN mkdir -p $JENKINS_HOME \
20      && chown ${uid}:${gid} $JENKINS_HOME \
21      && groupadd -g ${gid} ${group} \
22      && useradd -d "$JENKINS_HOME" -u ${uid} -g ${gid} -m -s /bin/bash ${user}
```



<https://github.com/jenkinsci/docker/blob/master/Dockerfile>

Multistage Build



```
FROM maven:3.3-jdk-8 as builder
COPY . /build/
WORKDIR /build
RUN mvn clean install

FROM openjdk:8-jre
COPY --from=builder /build/target/demoapp.jar /opt/
EXPOSE 8080
ENTRYPOINT ["java", "-jar", "/opt/demoapp.jar"]
```

```
# docker build -t myapp .
```

```
...
```

```
# docker build -t myapp --target builder .
```

```
...
```



[sbeliakou/springboot_example](#)

Try Out Online Test

Practice Test – Docker Image creation

Docker build is run by _____

Docker daemon

Docker CLI

Correct !



<https://vitalflux.com/docker-certification-practice-questions-dockerfile/>

<https://djitz.com/certification/docker-certified-associate-test-review-questions-set-1-image-creation/>



Running Containers

- **Running in detached mode**
- **Exposing Ports**
- **Managing Restart Policy**
- **Changing Workspace Directory**
- **Changing Runtime User**
- **Providing Environment Variables**
- **Providing Labels**

Working with Containers

Lifecycle:

- o [`docker create`](#) - creates a container but does not start it.
- o [`docker rename`](#) - allows the container to be renamed.
- o [`docker run`](#) - creates and starts a container in one operation.
- o [`docker rm`](#) - deletes a container.
- o [`docker update`](#) - updates a container's resource limits.

Starting and Stopping:

- o [`docker start`](#) - starts a container so it is running.
- o [`docker stop`](#) - stops a running container.
- o [`docker restart`](#) - stops and starts a container.
- o [`docker pause`](#) - pauses a running container, "freezing" it in place.
- o [`docker unpause`](#) - will unpauses a running container.
- o [`docker wait`](#) - blocks until running container stops.
- o [`docker kill`](#) - sends a SIGKILL to a running container.
- o [`docker attach`](#) - will connect to a running container.

Working with Containers

Info:

- [`docker ps`](#) - shows running containers.
- [`docker logs`](#) - gets logs from container.
- [`docker inspect`](#) - looks at all the info on a container.
- [`docker events`](#) - gets events from container.
- [`docker port`](#) - shows public facing port of container.
- [`docker top`](#) - shows running processes in container.
- [`docker stats`](#) - shows containers' resource usage statistics.
- [`docker diff`](#) - shows changed files in the container's FS.

Import / Export:

- [`docker cp`](#) - copies files or folders between a container and the local filesystem.
- [`docker export`](#) - turns container filesystem into tarball archive stream to STDOUT.

Executing Commands:

- [`docker exec`](#) - to execute a command in container.

docker run reference

```
# docker run options IMAGE Non_Default_CMD CMD_Args
```

options:

- d Run container in background and print container ID
- P Publish all exposed ports to random ports
- p Publish a container's port(s) to the host
- restart Restart policy to apply when a container exits (default "no")
- i Keep STDIN open even if not attached
- t Allocate a pseudo-TTY
- rm Automatically remove the container when it exits
- v Bind mount a volume
- e Set environment variables
- label Set meta data on a container
- log-driver Logging driver for the container

- u Username or UID (format: <name|uid>[:<group|gid>])
- w Working directory inside the container
- entrypoint Overwrite the default ENTRYPOINT of the image

<https://docs.docker.com/engine/reference/run/>

Running the Container

```
# docker run 50a986f614d5 # myhttpd:1.0
```

```
^C
```

```
# docker run -d myhttpd:1.0
```

```
9f761335efe268e9a82c4828d8f4be67b5824eb3266e8ba311343a7da45c67ff
```

```
# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
9f761335efe2	50a986f614d5	"/bin/sh -c 'httpd -l -v'"	5 seconds ago	Up 4 seconds	80/tcp	trusting_kilby

```
# docker run -P -d myhttpd:1.0
```

```
74954ff14ec5e53ac9925bfd2873c654fe8978657764b4162ac494fc9afaab9f
```

```
# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
74954ff14ec5	myhttpd:1.0	"/bin/sh -c 'httpd -l -v'"	9 seconds ago	Up 18 seconds	0.0.0.0:32768->80/tcp	fervent_noyle
9f761335efe2	myhttpd:1.0	"/bin/sh -c 'httpd -l -v'"	3 minutes ago	Up 3 minutes	80/tcp	trusting_kilby

```
# curl localhost:32768 # or curl <<VM_external_ip>>:32768
```

```
my httpd container
```

```
# docker run -d -p 8081:80 --name h8081 myhttpd:1.0
```

```
fca7f4525bc618e7c503b73bfa680c055300e8b5c767d48e33669831e0bc5bec
```

```
# docker run -d -p 127.0.0.1:8082:80 --name h8082 myhttpd:1.0
```

```
014e5efa5ca90d9b7e50eebef3c7f020f08f0b5238f98420681ee348a4097829
```

```
# docker ps --format "table {{.Names}}\t{{.Image}}\t{{.ID}}\t{{.Ports}}" -n2
```

CONTAINER ID	IMAGE	PORTS
fca7f4525bc6	myhttpd:1.0	0.0.0.0:8081->80/tcp
014e5efa5ca9	myhttpd:1.0	127.0.0.1:8082->80/tcp



Running the Container: Restart Policy

```
# docker run -d --restart=always --name sleeper centos sleep 5  
6c3d24b3f89f13de92e710fcbb5b343b4cb81e7454ecc79445e94f0c5ba31a49
```

```
# docker ps --format "table {{.Names}}\t{{.Image}}\t{{.ID}}\t{{.RunningFor}}\t{{.Status}}" -n1  
NAMES      IMAGE      CONTAINER ID      CREATED      STATUS  
sleeper    centos    33675bff7f47    2 seconds ago    Up 1 second
```

```
# docker ps --format "table {{.Names}}\t{{.Image}}\t{{.ID}}\t{{.RunningFor}}\t{{.Status}}" -n1  
NAMES      IMAGE      CONTAINER ID      CREATED      STATUS  
sleeper    centos    33675bff7f47    11 seconds ago   Up 4 seconds
```

```
# docker ps --format "table {{.Names}}\t{{.Image}}\t{{.ID}}\t{{.RunningFor}}\t{{.Status}}" -n1  
NAMES      IMAGE      CONTAINER ID      CREATED      STATUS  
sleeper    centos    33675bff7f47    About a minute ago  Restarting (0) 6 seconds ago
```

```
# docker inspect -f "{{ .RestartCount }}" sleeper  
19
```

Policy	Result
no	Do not automatically restart the container when it exits. This is the default.
on-failure[:max-retries]	Restart only if the container exits with a non-zero exit status. Optionally, limit the number of restart retries the Docker daemon attempts.
always	Always restart the container regardless of the exit status. When you specify always, the Docker daemon will try to restart the container indefinitely. The container will also always start on daemon startup, regardless of the current state of the container.
unless-stopped	Always restart the container regardless of the exit status, including on daemon startup, except if the container was put into a stopped state before the Docker daemon was stopped.



More details

Running Containers in Interactive Mode

```
[root@localhost ~]# docker run centos cat /etc/redhat-release  
CentOS Linux release 7.5.1804 (Core)
```

```
[root@localhost ~]# docker run ubuntu cat /etc/lsb-release  
DISTRIB_ID=Ubuntu  
DISTRIB_RELEASE=18.04  
DISTRIB_CODENAME=bionic  
DISTRIB_DESCRIPTION="Ubuntu 18.04.1 LTS"
```

```
[root@localhost ~]# docker run -it centos bash  
[root@dfc1b0d4f6a5 /]# cat /etc/redhat-release  
CentOS Linux release 7.5.1804 (Core)  
[root@dfc1b0d4f6a5 /]# yum install curl wget # and so on  
...  
[root@dfc1b0d4f6a5 /]# exit  
exit  
[root@localhost ~]#
```

```
[root@localhost ~]# docker ps --format "table {{.ID}}\t{{.RunningFor}}\t{{.Status}}" -n1 -a  
CONTAINER ID        CREATED          STATUS  
dfc1b0d4f6a5        7 minutes ago   Exited (1) 1 minutes ago
```

```
[root@localhost ~]# docker rm dfc1b0d4f6a5  
dfc1b0d4f6a5
```

```
[root@localhost ~]# docker run --rm centos cat /etc/redhat-release  
CentOS Linux release 7.5.1804 (Core)
```

Executing Commands Inside Running Container

```
# docker run -d centos sleep infinity  
9626a94669c935ea140bcb9ea83339bd325b28195fc088a690b620eb12902b33
```

```
# docker ps -l  
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES  
9626a94669c9 centos "sleep infinity" 11 seconds ago Up 10 seconds priceless_einstein
```

```
# docker exec -it 9626a94669c9 bash  
[root@9626a94669c9 /]# ps -ef  
UID      PID  PPID  C STIME TTY          TIME CMD  
root        1      0  0 13:43 ?        00:00:00 sleep infinity  
root       34      0  2 13:48 pts/0    00:00:00 bash  
root       47     34  0 13:48 pts/0    00:00:00 ps -ef  
[root@9626a94669c9 /]# exit  
exit  
#
```

Stopping/Deleting Containers

```
# docker ps --format "table {{.Image}}\t{{.Names}}\t{{.ID}}\t{{.RunningFor}}\t{{.Status}}"
IMAGE      NAMES      CONTAINER ID      CREATED      STATUS
myhttpd:1.0 h8081     fca7f4525bc6    About an hour ago Up About an hour
myhttpd:1.0 h8082     014e5efa5ca9    About an hour ago Up About an hour
```

```
# docker stop h8082
014e5efa5ca9
```

```
# docker rm 014e5efa5ca9
014e5efa5ca9
```

```
# docker rm $(docker stop 014e5efa5ca9)
014e5efa5ca9
```

```
# docker rm $(docker stop $(docker ps -a -q))
014e5efa5ca9
fca7f4525bc6
```



```
# docker container prune
WARNING! This will remove all stopped containers.
Are you sure you want to continue? [y/N]
```



```
# docker image prune
WARNING! This will remove all dangling images.
Are you sure you want to continue? [y/N]
you want to continue? [y/N]
```

Changing Container's Build Defaults

1. Default User

```
# docker run jenkins id  
uid=1000(jenkins) gid=1000(jenkins) groups=1000(jenkins)
```

```
# docker run --user 0 jenkins id  
uid=0(root) gid=0(root) groups=0(root)
```

--user == -u

```
# docker run --user 1000:0 jenkins id  
uid=1000(jenkins) gid=0(root) groups=0(root)
```

```
# docker run --group-add 123 jenkins id  
uid=1000(jenkins) gid=1000(jenkins) groups=1000(jenkins),123
```



```
...  
RUN useradd jenkins -u 1000  
USER jenkins
```

2. Default Workdir

```
# docker run jenkins pwd  
/
```

```
# docker run --workdir /var/jenkins_home jenkins pwd  
/var/jenkins_home
```

--workdir == -w

```
# docker run -w $(pwd) -v $(pwd):$(pwd) maven clean package
```



```
...  
WORKDIR /
```

Changing Container's Build Defaults

3. Default Entrypoint

```
# docker run -it myhttpd:1.0 bash  
# echo $?  
1
```



```
...  
ENTRYPOINT ["httpd"]  
...
```

```
# docker run -it --entrypoint=/bin/bash myhttpd:1.0  
[root@e4e7976cf838 /]# ps -ef  
UID      PID  PPID  C STIME TTY          TIME CMD  
root        1      0  1 19:27 pts/0    00:00:00 /bin/bash  
root       14      1  0 19:27 pts/0    00:00:00 ps -ef  
[root@e4e7976cf838 /]# exit  
#
```

4. Environment Variables

```
# docker run -it -e MYVAR="My Variable" centos env | grep MYVAR  
MYVAR=My Variable
```

```
# docker run -it --env-file <(env| grep ARM | cut -f1 -d=) centos env | grep ARM  
ARM_SUBSCRIPTION_ID=64a3f30f-xxx-xxxx-xxxx-yyyyfe63fe9a  
ARM_TENANT_ID=bd5c6713-xxx-yyyy-xxxx-78f2d078e543  
ARM_CLIENT_SECRET=xxxxxxxxxxxxxxxx  
ARM_CLIENT_ID=808f38ed-xxxx-xxxx-yyyy-ebbce91bcfee
```

Running Containers with Labels

```
# docker run -d --label app=web1 nginx
```

```
# docker run -d --label app=web2 nginx
```

```
# docker run -d --label app=web3 nginx
```

```
# docker ps --filter 'label=app=web1' --format "table {{.ID}}\t{{.Status}}"
```

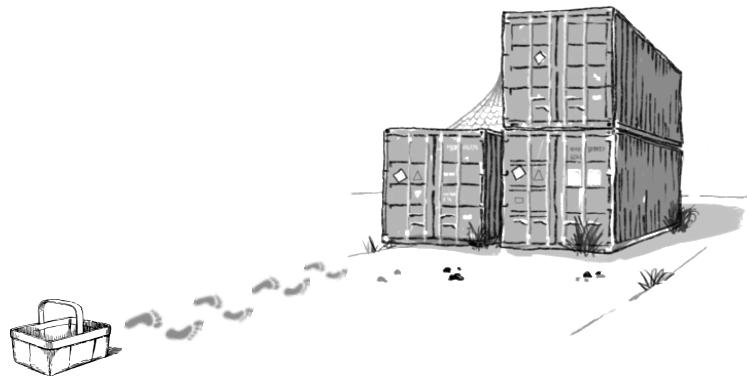
CONTAINER ID	STATUS
978dc141e8e0	Up 2 minutes

```
# docker stop $(docker ps --filter 'label=app=web1' -q)
```

```
978dc141e8e0
```

```
# docker ps --filter 'label=app=web1' --format "table {{.ID}}\t{{.Status}}"
```

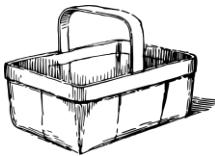
CONTAINER ID	STATUS
<< no containers, all stopped by command above >>	



Mounting Data from the Host

- **Mounting Folders**
- **Mounting Files**
- **R/O Mounts**
- **Mounting TMPFS**

Use Cases



1. Running Containers with Local Configs / Data

2. Running Containers for Processing Data (Tools)

3. Saving Changed Data (Stateful Application)

4. Sharing Data between Containers

5. Mounting Application Logs to the Host ***

Mounting Data from Host

```
docker run ... -v ${PATH_HOST}:${PATH_CONTAINER} <<image_name>>
docker run ... -v ${PATH_HOST}:${PATH_CONTAINER}:ro <<image_name>>
```

Examples:

Run httpd service with custom "index.html"

```
# ls -l .
total 40
-rw-r--r-- 1 sbeliakou wheel 119 July 25 22:05 index.html
```

```
# docker run -d -P -v $(pwd):/var/www/html httpd
# docker run -d -P -v $(pwd):/var/www/html:ro httpd
```

```
# docker run -d -p 127.0.0.1:8080:80 -v $(pwd):/var/www/html:ro httpd
```

Run nginx service with custom config

```
# docker run -d -v $(pwd)/nginx.conf:/etc/nginx/nginx.conf nginx
```

Use Cases

```
# docker run --rm -v $(pwd):$(pwd) -w $(pwd) maven:3.3-jdk-8 clean package
```

```
# docker run -d -v /var/log/httpd:/var/log/httpd httpd
```

```
# docker run -d -v /var/log/tomcat:/usr/local/tomcat/logs tomcat
```

```
# docker run -d -v /data:/etc/mongo mongo
```



https://hub.docker.com/_/jenkins/

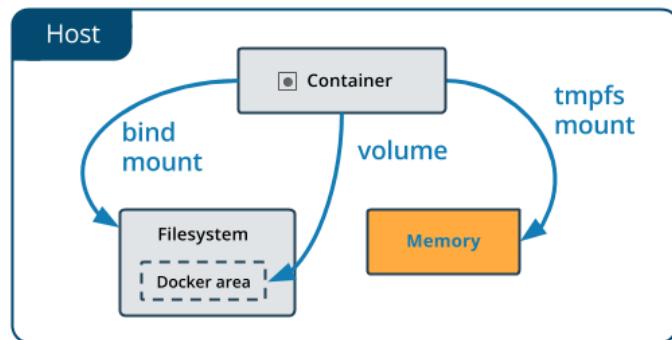
```
# docker run -p 8080:8080 -p 50000:50000 -v /your/dir:/var/jenkins_home jenkins
```

Mounting tmpfs

A *tmpfs* mount is temporary, and only persisted in the host memory. When the container stops, the tmpfs mount is removed, and files written there won't be persisted.

Limitations of *tmpfs* mounts:

- You can't share tmpfs mounts between containers.
- This functionality is only available if you're running Docker on Linux.



```
# docker run -d \
--mount type=tmpfs,destination=/app \
nginx:latest
```

```
# docker run -d \
--tmpfs /app \
nginx:latest
```

<https://docs.docker.com/storage/tmpfs/#use-a-tmpfs-mount-in-a-container>



to be continued

Thank you for your attention!

Siarhei Beliakou,

2019