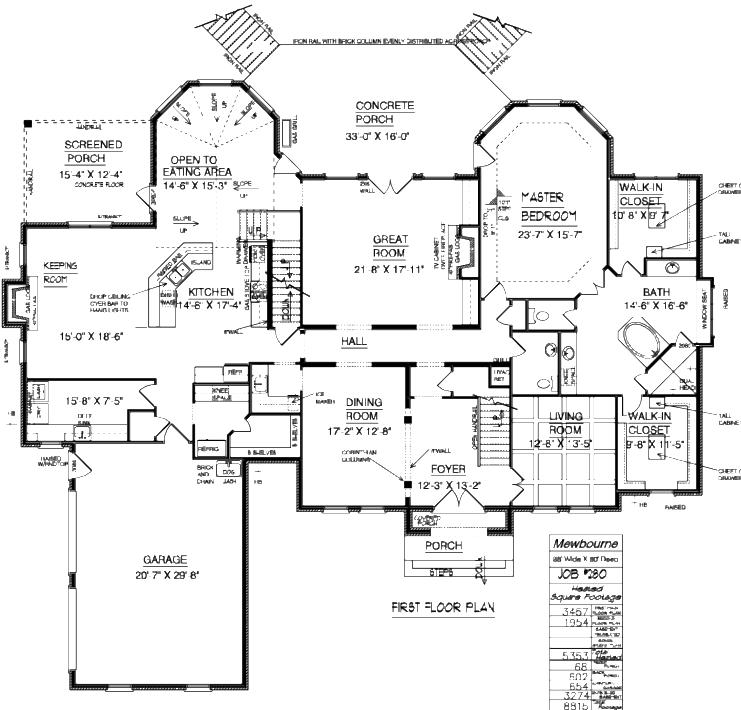




# AGENDA

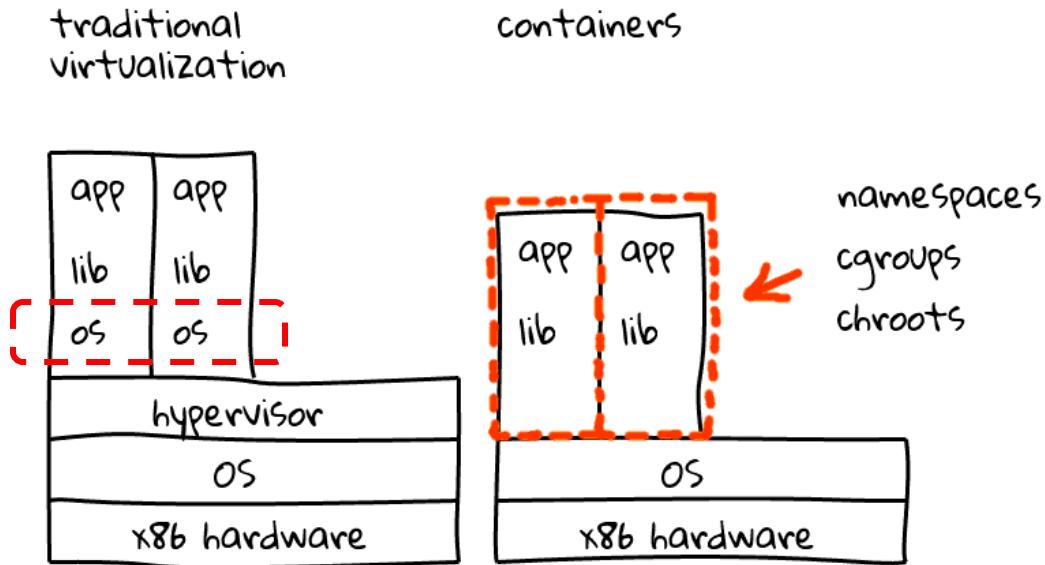
- **Linux Kernel Namespaces and Control Groups**
- **Linux Process Capabilities**
- **Docker-In-Docker**
- **Remote API**
- **Docker Plugins**
- **Monitoring Containers**
- **Containers Security**



# Linux Kernel Namespaces and CGroups

- **Namespaces**
    - pid
    - net
    - uts
  - **Cgroups**
    - Memory Constraints
    - CPU Constraints

# What is a Container?



## Each Container has:

- Own process space
- Own network interface
- Can run as root
- Can install packages
- Can run services
- etc

## Why Containers run unlike VMs:

- It uses host's Kernel
- It doesn't need INIT system
- It doesn't need syslog, cron and so on

# Cgroups and Namespaces

---

**Cgroups** = limits how much you can use

- Resource metering and limiting:
  - memory
  - CPU
  - block I/O
  - network
- Device node (/dev/\*) access control

**Namespaces** = limits what you can see (and therefore use)

- Provide processes with their own view of the system
- Multiple namespaces:
  - pid – isolates the process ID number space
  - net – manages network devices
  - mnt – to see distinct single- directory hierarchies
  - uts – isolating hostnames
  - ipc – manages shared memory areas, message queues, and semaphores
- Each process is in one namespace of each type

# Namespace PID

Run Container in Host PID Namespace:

```
# sleep 1000 &
[1] 19005
$ ps -ef | grep sleep | grep -v grep
root      19005      1  0 21:57 pts/0    00:00:00 sleep 1000
$ docker run --pid=host centos kill -9 19005
$ ps -ef | grep sleep
```

Run Container in another container PID Namespace:

```
# docker run -d --name=web-container nginx
# docker run --rm --pid=container:web-conatiner centos ps
PID TTY          TIME CMD
 1 ?        00:00:00 nginx
 12 ?       00:00:00 ps
```

# Namespace NET

```
# docker run -d --name nginx nginx
# docker run -d -it --name net-tools --net=container:nginx sbeliakou/net-tools
```

```
# docker exec nginx hostname -i
172.17.0.2
# docker exec net-tools hostname -i
172.17.0.2
```

```
# docker exec net-tools ps -ef
PID  USER      TIME  COMMAND
 1  root      0:00 /bin/bash
 11 root      0:00 ps -ef
```

```
# docker exec net-tools netstat -natpl
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address          Foreign Address        State      PID/Program name
tcp        0      0 0.0.0.0:80              0.0.0.0:*            LISTEN     -
```

```
# docker exec net-tools curl -sIL localhost | egrep "(HTTP|Server)"
HTTP/1.1 200 OK
Server: nginx/1.15.2
```

```
# curl -sIL 172.17.0.2 | egrep "(HTTP|Server)"
HTTP/1.1 200 OK
Server: nginx/1.15.2
```

# Namespace UTS (System Name and Domain)

By default, a container runs with a UTS namespace (which defines the system name and domain) that is different from the UTS namespace of the host. To make a container use the same UTS namespace as the host, you can use the `--uts=host` option with the docker create and docker run commands from version 1.7.0 of Docker onward. This setting allows the container to track the UTS namespace of the host or to set the host name and domain from the container.

```
[root@docker-host ~]# hostname  
docker-host
```

```
[root@docker-host ~]# docker run --rm busybox hostname  
0feabf511a54
```

```
[root@docker-host ~]# docker run --rm --uts=host busybox hostname  
docker-host
```

# Cgroups: Memory Limits

```
# docker run -d centos sleep infinity  
60614832f3ba25962c584d209a0ba1f99335c3c980563c72715ac14175fada5f
```

```
# docker stats --no-stream --format "table {{.Container}}\t{{.MemUsage}}" 60614832f3ba  
CONTAINER           MEM USAGE / LIMIT  
60614832f3ba      92KiB / 991.6MiB
```

```
# docker run -d -m 300M centos sleep infinity  
b133cd44aa414645f94474ffdc49b7d31b7cccfcb14a122f8f0700bb0b7c80bd
```

```
# docker stats --no-stream --format "table {{.Container}}\t{{.MemUsage}}" b133cd44aa41  
CONTAINER           MEM USAGE / LIMIT  
b133cd44aa41       1016KiB / 300MiB
```

```
# docker run -d -m 300M --memory-reservation 100M centos sleep infinity  
af6df26f67538433932f2ab9957afa2eaec586dedae5a99c837df2a5be6e2c18
```

```
# docker stats --no-stream --format "table {{.Container}}\t{{.MemUsage}}"  
CONTAINER           MEM USAGE / LIMIT  
af6df26f6753       968KiB / 300MiB  
b133cd44aa41       1016KiB / 300MiB  
60614832f3ba       92KiB / 991.6MiB
```

Check out documentation:  
container memory [constraints](#)  
docker stats [reference](#)

# Cgroups: CPU Limits

CPU limits are based on shares. These shares are a weight between how much processing time one process should get compared to another. If a CPU is idle, then the process will use all the available resources. If a second process requires the CPU then the available CPU time will be shared based on the weighting.

- |               |   |
|---------------|---|
| --cpu-shares  | - CPU shares (relative weight).                     |
| --cpu-period  | - Limit CPU CFS (Completely Fair Scheduler) period. |
| --cpu-quota   | - Limit CPU CFS (Completely Fair Scheduler) quota.  |
| --cpuset-cpus | - CPUs in which to allow execution (0-3, 0,1).      |
| --cpuset-mems | - MEMs in which to allow execution (0-3, 0,1).      |

The **--cpu-quota** option specifies the number of microseconds that a container has access to CPU resources during a period specified by **--cpu-period**. As the default value of **--cpu-period** is 100000, setting the value of **--cpu-quota** to 25000 limits a container to 25% of the CPU resources. By default, a container can use all available CPU resources, which corresponds to a **--cpu-quota** value of -1

```
# docker run -it -d --cpu-quota=25000 --name cpu0.25 benhall/stress
```

```
# docker stats --no-stream --format "table {{.Name}}\t{{.CPUPerc}}" cpu0.25
NAME          CPU %
cpu0.25      25.03%
```



## Linux Process Capabilities

- Capabilities Overview
- Enabling Capabilities
- SYS\_TIME Example
- NET\_ADMIN Example
- SYS\_ADMIN Example
- Best Practices (Sources)

# Linux Process Capabilities

Unix was designed to run two categories of processes :

- privileged processes, with the user id 0 which is the root user.
- unprivileged processes where the user id is different from zero.

While privileged processes bypass all kernel permission checks, all the other processes are subject to a permission check ( the Unix-like system Kernel make the permission check based on the UID and the GID of the non root process.

Starting with Kernel 2.2, Linux divides the privileges into distinct units that we call capabilities. Capabilities can be independently enabled and disabled, which is a great security and permission management feature. A process can have CAP\_NET\_RAW capability while it is denied to use the CAP\_NET\_ADMIN capability. This is just an example explaining how capabilities works. You can find the complete list of the other capabilities in Linux man-page:

```
$ man capabilities
```

<http://man7.org/linux/man-pages/man7/capabilities.7.html>

# Managing Capabilities on a Container

```
# docker run --help
--cap-add: Add Linux capabilities
--cap-drop: Drop Linux capabilities
--privileged=false: Give extended privileges to this container
--device=[]: Allows you to run devices inside the container without the privileged flag.
```

```
# docker run --privileged=true ...
```

## List of Capabilities:

<https://docs.docker.com/engine/reference/run/#runtime-privilege-and-linux-capabilities>

## Docker Security Article:

<https://linux-audit.com/docker-security-best-practices-for-your-vessel-and-containers/>

# Capability: SYS\_TIME

```
[root@localhost ~]# timedatectl | grep NTP
    NTP enabled: yes
    NTP synchronized: yes
```

```
[root@localhost ~]# date
Tue  7 Aug 21:54:57 UTC 201
```

```
[vagrant@localhost ~]# docker run --rm centos date
Tue Aug  7 21:55:06 UTC 2018
```

```
[vagrant@localhost ~]# docker run --rm --cap-add=SYS_TIME centos date -s '+2 hours'
Tue Aug  7 23:55:15 UTC 2018
```

```
[vagrant@localhost ~]# date
Tue  7 Aug 23:55:17 UTC 2018
```

```
[root@localhost ~]# timedatectl | grep NTP
    NTP enabled: yes
    NTP synchronized: no
```

# Capability: NET\_ADMIN

**Perform various network-related operations:**

- interface configuration;
- administration of IP firewall, masquerading, and accounting
- modify routing tables;
- bind to any address for transparent proxying;
- set type-of-service (TOS)
- clear driver statistics;
- set promiscuous mode;
- enabling multicasting;

```
$ docker run --cap-add=NET_ADMIN ubuntu:14.04 ip link add dummy0 type dummy
```

```
$ ip link | grep -A1 dummy
```

```
31: dummy0: <BROADCAST,NOARP> mtu 1500 qdisc noop state DOWN mode DEFAULT group default qlen 1000  
link/ether 7e:32:dd:ec:f1:22 brd ff:ff:ff:ff:ff:ff
```

```
docker run --cap-add=NET_ADMIN ubuntu:14.04 sh -c 'ip link eth0 down'
```

# Capability: SYS\_ADMIN

```
$ hostname  
docker-host
```

```
$ docker run --rm busybox hostname  
0feabf511a54
```

```
$ docker run --rm --uts=host busybox hostname  
docker-host
```

```
$ docker run --rm --uts=host busybox hostname HOST  
hostname: sethostname: Operation not permitted
```

```
$ docker run --rm --uts=host busybox hostname HOST  
hostname: sethostname: Operation not permitted
```

```
$ docker run --rm --uts=host --cap-add=SYS_ADMIN busybox hostname HOST  
$ hostname  
HOST
```

# Lean VM

```
$ docker run -d centos /usr/sbin/init
$ docker exec -it $(docker ps -ql) bash
[root@b81e02718292 /]# systemctl status
Failed to get D-Bus connection: Operation not permitted
```

```
$ docker run -d -v /sys/fs/cgroup centos /usr/sbin/init
$ docker exec -it $(docker ps -ql) bash
[root@2a3feeafc6fb /]# systemctl status
Failed to get D-Bus connection: Operation not permitted
```

```
$ docker run -d -v /sys/fs/cgroup --cap-add=SYS_ADMIN centos /usr/sbin/init
$ docker exec -it $(docker ps -ql) bash
[root@7033599262f5 /]# systemctl status
● 7033599262f5
    State: running
    Jobs: 0 queued
    Failed: 0 units
    Since: Mon 2018-10-15 18:39:54 UTC; 6s ago
    CGroup: /docker/7033599262f54477c13ccd4d2ebf3bb8af52eec5d41295c973b2a022b1de0ff9
            |- 1 /usr/sbin/init
            |- 31 bash
            |- 44 systemctl status
            |- 45 systemctl status
            |- system.slice
                ...

```

# Lean VM



```
FROM centos
```

[CentOS.SSHD.Dockerfile](#)

```
RUN yum install -y systemd openssh-server sudo && systemctl enable sshd  
RUN useradd devops && echo devops | passwd devops --stdin
```

```
EXPOSE 22
```

```
ENTRYPOINT ["/usr/sbin/init"]
```



```
version: "2.3"
```

[docker-compose.yml](#)

```
services:
```

```
centosvm:
```

```
build:
```

```
context: .
```

```
dockerfile: CentOS.SSHD.Dockerfile
```

```
volumes:
```

```
- /sys/fs/cgroup
```

```
cap_add:
```

```
- SYS_ADMIN
```

# Lean VM

```
$ docker-compose up -d
Creating network "leanvm_default" with the default driver
Creating leanvm_centosvm_1 ... done
```

```
$ docker-compose ps
      Name           Command       State    Ports
----- 
leanvm_centosvm_1   /usr/sbin/init   Up        22/tcp
```

```
$ docker-compose up -d --scale centosvm=5
Starting leanvm_centosvm_1 ... done
Creating leanvm_centosvm_2 ... done
Creating leanvm_centosvm_3 ... done
Creating leanvm_centosvm_4 ... done
Creating leanvm_centosvm_5 ... done
```

```
$ docker-compose ps
      Name           Command       State    Ports
----- 
leanvm_centosvm_1   /usr/sbin/init   Up        22/tcp
leanvm_centosvm_2   /usr/sbin/init   Up        22/tcp
leanvm_centosvm_3   /usr/sbin/init   Up        22/tcp
leanvm_centosvm_4   /usr/sbin/init   Up        22/tcp
leanvm_centosvm_5   /usr/sbin/init   Up        22/tcp
```

```
# docker inspect \
  -f '{{range .NetworkSettings.Networks}}{{.IPAddress}}{{end}}' \
$(docker-compose ps -q)
172.24.0.2
172.24.0.5
172.24.0.6
172.24.0.4
172.24.0.3
```

```
# ssh devops@172.25.0.4
The authenticity of host '172.25.0.4 (172.25.0.4)' can't be established.
ECDSA key fingerprint is SHA256:WRIJRV4kwjpx34t3HnHP2RdYsXVhd4j0EAALKw0IddQ.
ECDSA key fingerprint is MD5:42:87:a3:65:1b:b8:64:37:7a:9f:6f:b9:66:22:f3:f1.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '172.25.0.4' (ECDSA) to the list of known hosts.
devops@172.25.0.4's password:
[devops@83c5e3cd314e ~]$
```

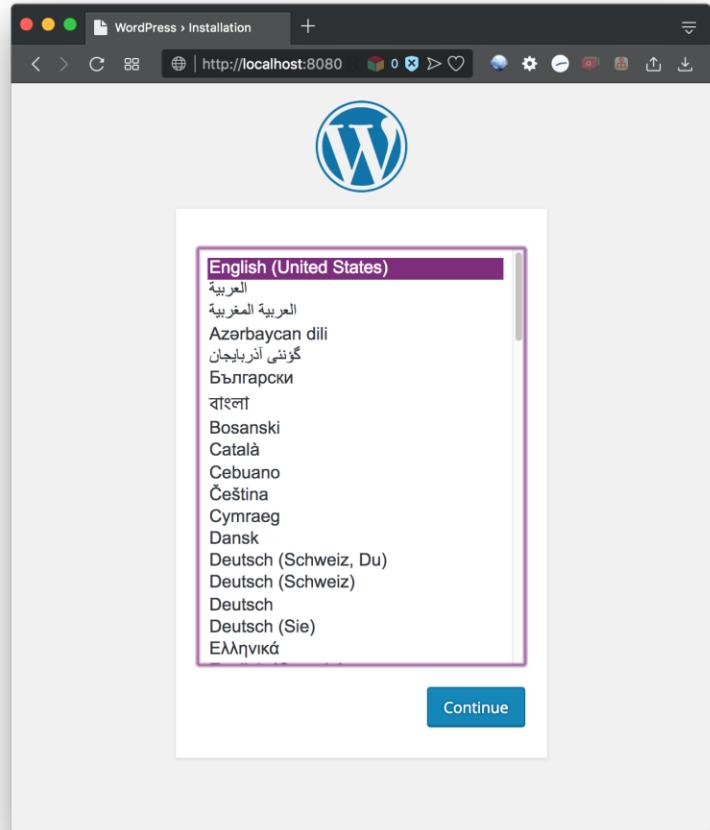
# WordPress Example



```
version: "3"

services:
  wordpress:
    image: endophage/wordpress:lean
    links:
      - mysql
  ports:
    - "8080:80"
  environment:
    - DB_PASSWORD=2671d40f658f595
  cap_drop:
    - ALL
  cap_add:
    - SETUID
    - SETGID
    - DAC_OVERRIDE
    - NET_BIND_SERVICE

mysql:
  image: mariadb:10.1.10
  environment:
    - MYSQL_DATABASE=wordpress
    - MYSQL_ROOT_PASSWORD=2671d40f658f595
  ports:
    - "3306"
```



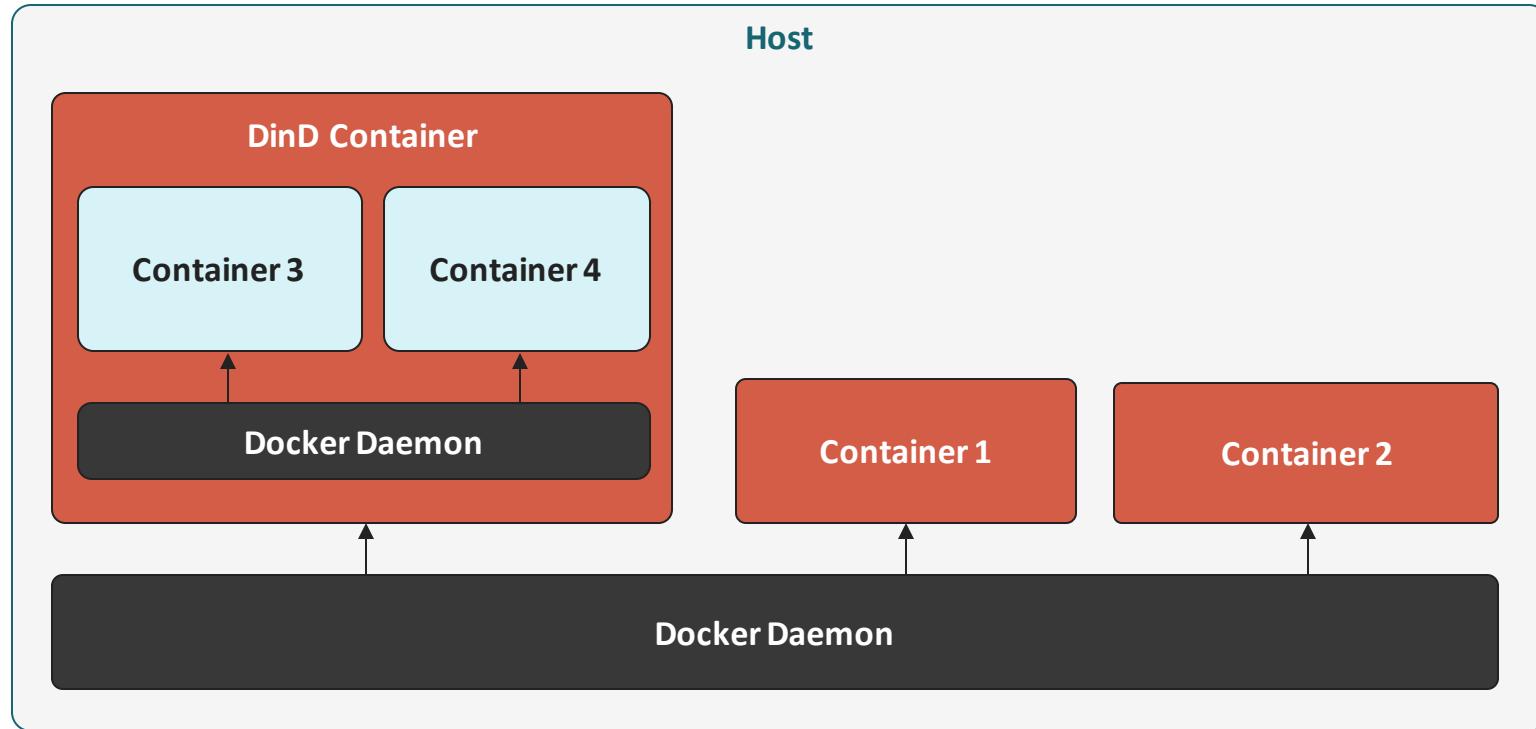


## Docker-in-Docker

- How does this work?
- Do you really need this?

# Docker in Docker (DiD)

How it looks:



# Docker in Docker (DinD)

Run "dind" container:

```
# docker run --privileged -v /usr/local/bin -v /var/run --name dind -d docker:dind
```

Run "nginx" in "dind" container (from inside "dind"):

```
# docker exec dind docker run -d --name=nginx1 nginx
```

Run "nginx" in "dind" container (from outside "dind"):

```
# docker run --volumes-from dind -e PATH=/usr/local/bin busybox \
docker run --name=nginx2 -d nginx
```

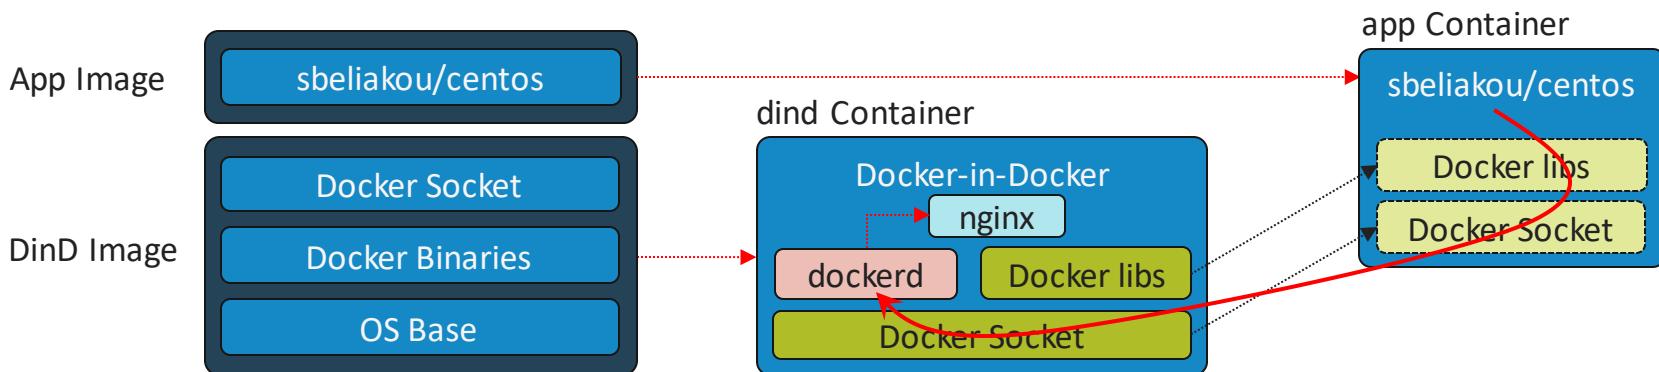
How goes?

```
# docker ps --format "table {{.Names}}\t{{.Status}}"
NAMES          STATUS
dind          Up 3 minutes
# docker exec -it dind docker ps --format "table {{.Names}}\t{{.Status}}"
NAMES          STATUS
nginx2         Up About a minute
nginx1         Up 3 minutes
```

# Understanding DinD

```
$ docker run --name dind -d \
  --privileged \
  -v /usr/local/bin \
  -v /var/run \
  docker:dind
```

```
$ docker run --rm -ti --volumes-from=dind sbeliakou/centos bash
bash-4.2# docker run -d -P nginx
bash-4.2# docker ps
```



# Data Saving Scenario

We need to create a volume for storing our dind's containers:

```
# docker volume create dind_data
```

Run our dind container with our volume:

```
# docker run --privileged \
-v /usr/local/bin \
-v /var/run \
-v dind_data:/var/lib/docker \
--name dind -d docker:dind
```

Let's start Nginx Container inside DinD:

```
# docker exec dind docker run -d --name=nginx --restart=always nginx
```

Stopping and Removing DinD:

```
# docker rm $(docker stop dind)
```

Start DinD again:

```
# docker run --privileged ... # command as above
```



## Docker Remote API

- **Configuring Remote Access**
- **Accessing Remote Docker Daemon**
- **Securing Remote Access**

# Configuring Docker Daemon for Remote Access

## ! systemd vs daemon.json

Configuring Docker to listen for connections using both the *systemd* unit file and the *daemon.json* file causes a conflict that prevents Docker from starting.

/etc/docker/daemon.json

```
{  
  "hosts": [  
    "unix:///var/run/docker.sock",  
    "tcp://127.0.0.1:2375",  
    "192.168.56.15"  
  ]  
}
```

[More Configuration Options](#)

/usr/lib/systemd/system/docker.service

```
...  
[Service]  
Type=notify  
ExecStart=/usr/bin/dockerd -H unix:///  
ExecReload=/bin/kill -s HUP $MAINPID  
TimeoutSec=0  
RestartSec=2  
Restart=always  
...
```

```
$ systemctl daemon-reload  
$ systemctl restart docker.service
```

```
# curl 192.168.56.15:2375/_ping  
OK
```

# Accessing Remote Docker Daemon

```
# export DOCKER_HOST=tcp://192.168.56.15:2375  
# docker ps  
...
```

```
# docker -H tcp://192.168.56.15:2375 ps  
...
```

```
# curl 192.168.56.15:2375/info
```

```
# docker info
```

```
# curl 192.168.56.15:2375/images/json
```

```
# docker images
```

```
# curl 192.168.56.15:2375/containers/json
```

```
# docker ps
```

```
# curl 192.168.56.15:2375/containers/json?all=1
```

```
# docker ps -a
```

# Securing Remote Access

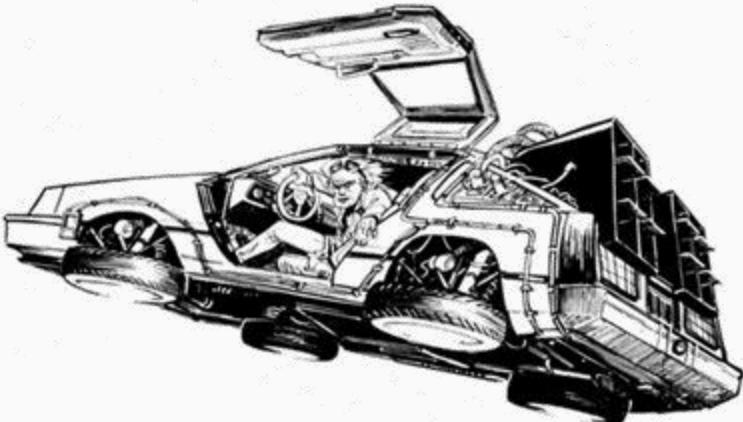
---

<https://docs.docker.com/engine/security/https/>

[https://docs.docker.com/engine/extend/plugins\\_authorization/](https://docs.docker.com/engine/extend/plugins_authorization/)



1. Enable TLS (tlsverify, tlscacert)
2. Enable 2 way Authentication (with AuthZ plugins)



## Docker Plugins

- Local-Persist
- SshFS

# Docker Engine Plugins

You can extend the capabilities of the Docker Engine by loading third-party plugins.

Currently Docker supports following driver plugins:

- authorization
- volume
- Network

For example, let's install [Local Persist Plugin](#):

```
$ curl -fsSL https://raw.githubusercontent.com/CWSpear/local-persist/master/scripts/install.sh | sudo bash
```

Creating "jenkins" Volume with "local-persist" driver:

```
$ docker volume create -d local-persist -o mountpoint=/data/jenkins --name=jenkins
```

Run Jenkins with the volume created above:

```
$ docker run -d -v jenkins:/var/jenkins_home jenkins/jenkins:2.146
```

# Using Local-Persist Volume with Docker-Compose



```
version: '2'

services:
  jenkins:
    image: jenkins/jenkins:2.146
    ports:
      - 80:8080
    volumes:
      - jenkins:/var/jenkins_home

volumes:
  jenkins:
    driver: local-persist
    driver_opts:
      mountpoint: /data/jenkins
```

# Docker volume plugin for sshFS

Install plugin with :

<https://github.com/vieux/docker-volume-sshfs>

```
# docker plugin install vieux/sshfs
```

Plugin "vieux/sshfs" is requesting the following privileges:

- network: [host]
- mount: [/var/lib/docker/plugins/]
- mount: []
- device: [/dev/fuse]
- capabilities: [CAP\_SYS\_ADMIN]

Do you grant the above permissions? [y/N] y

latest: Pulling from vieux/sshfs

52d435ada6a4: Download complete

Digest: sha256:1d3c3e42c12138da5ef7873b97f7f32cf99fb6edde75fa4f0bcf9ed277855811

Status: Downloaded newer image for vieux/sshfs:latest

Installed plugin vieux/sshfs

Check the Plugin Status:

```
$ docker plugin ls | grep sshfs
```

0c3c6b601473	vieux/sshfs:latest	sshFS plugin for Docker	true

# Docker volume plugin for sshFS

Create the Volume:

```
# docker volume create -d vieux/sshfs \
    -o sshcmd=<user@host:path> \
    -o password=<password> \
    [-o port=<port>] [-o <any_sshfs_-o_option> ] \
    sshvolume sshvolume
```

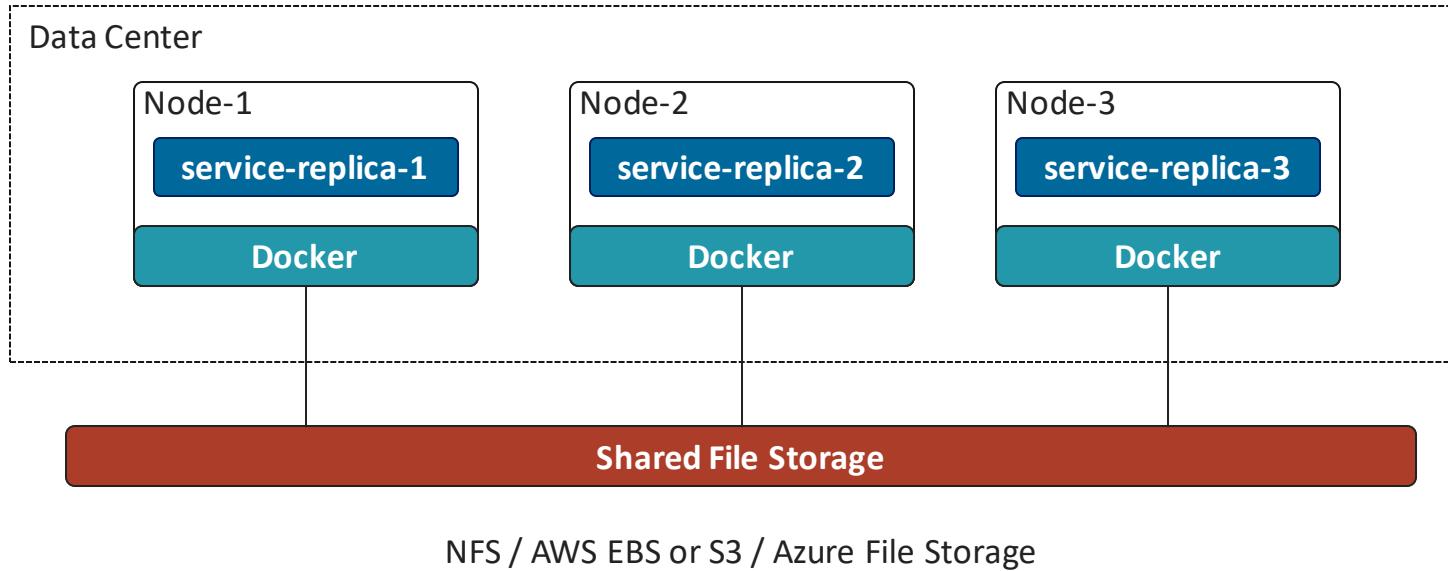
Check the volume :

```
# docker volume ls | grep sshvolume
vieux/sshfs      sshvolume
```

Use the volume :

```
# docker run -it -v sshvolume:<path_in_container> busybox ls <path>
```

# When is It Useful?



# Using Volume Drivers in Clouds



<https://docs.aws.amazon.com/AmazonECS/latest/developerguide/docker-volumes.html>



<https://docs.docker.com/docker-for-azure/persistent-data-volumes/>

Example for Azure:

```
# docker plugin install \
--alias cloudstor:azure \
--grant-all-permissions docker4x/cloudstor:18.03.1-ce-azure1 \
CLOUD_PLATFORM=AZURE \
AZURE_STORAGE_ACCOUNT_KEY=$AZURE_STORAGE_ACCOUNT_KEY \
AZURE_STORAGE_ACCOUNT=$AZURE_STORAGE_ACCOUNT
```

# Rexray

Provider	Storage Platform	Docker
Amazon EC2	<a href="#">EBS</a>	✓
	<a href="#">EFS</a>	✓
	<a href="#">S3FS</a>	✓
Ceph	<a href="#">RBD</a>	✓
Dell EMC	<a href="#">Isilon</a>	✓
	<a href="#">ScaleIO</a>	✓
DigitalOcean	<a href="#">Block Storage</a>	✓
FittedCloud	<a href="#">EBS Optimizer</a>	✓
Google	<a href="#">GCE Persistent Disk</a>	✓
Microsoft	<a href="#">Azure Unmanaged Disk</a>	✓
OpenStack	<a href="#">Cinder</a>	✓
VirtualBox	<a href="#">Virtual Media</a>	✓

REX-Ray provides a vendor agnostic storage orchestration engine. The primary design goal is to provide persistent storage for Docker, Kubernetes, and Mesos.

```
docker plugin install \
  rexray/ebs \
  EBS_ACCESSKEY=access_key \
  EBS_SECRETKEY=secret_key
```

```
docker run -ti \
  --volume-driver=rexray/ebs \
  --volume test:/test \
  busybox
```

<https://github.com/rexray/rexray>



## Monitoring Containers

- **What you should look after**
- **Containers Monitoring with cAdvisor**
- **Collect Docker metrics with Prometheus**

# What You Should Look after

---

At least:

- |                          |  |
|--------------------------|--|
| <b>Image Count</b>       | - The count of active and inactive images stored locally by Docker.  |
| <b>Image Size</b>        | - The total disk usage of active and reclaimable images stored locally.  |
| <b>Container Count</b>   | - The count of containers, segregated by state: created, running, paused, restarting, removing, exited and dead. |
| <b>Container Size</b>    | - The total size of disk used by containers, with used and reclaimable parts tracked separately.                 |
| <b>Volume Count</b>      | - The count of active and inactive volumes maintained by this Docker daemon.                                     |
| <b>Volume Size</b>       | - The total size of disk used by volumes, with used and reclaimable parts tracked separately.                    |
| <b>CPU Usage</b>         | - The percentage of host CPU used by all currently running containers.   |
| <b>Memory Usage</b>      | - The amount of host memory used by all currently running containers.  |
| <b>Network Bandwidth</b> | - The transmit and receive bandwidth of host network used by all currently running containers.                   |
| <b>Disk Throughput</b>   | - The disk read and write throughput (bytes/second) used by all currently running containers.                    |
| <b>Process Count</b>     | - The total number of processes running across all containers currently.   |

# Monitoring Containers with cAdvisor



```
version: "3.6"

services:
  web:
    image: google/cadvisor:latest
    ports:
      - 0.0.0.0:9104:8080
    volumes:
      - /:/rootfs:ro
      - /var/run:/var/run:rw
      - /sys:/sys:ro
      - /var/lib/docker/:/var/lib/docker:ro
      - /dev/disk/:/dev/disk:ro
```

```
# docker-compose up -d
```

cAdvisor - / Not Secure | 192.168.56.15:9104/containers/

## Memory

Reservation unlimited

Limit 487.58 MB

Swap Limit 2.00 GB

## Usage

### Overview

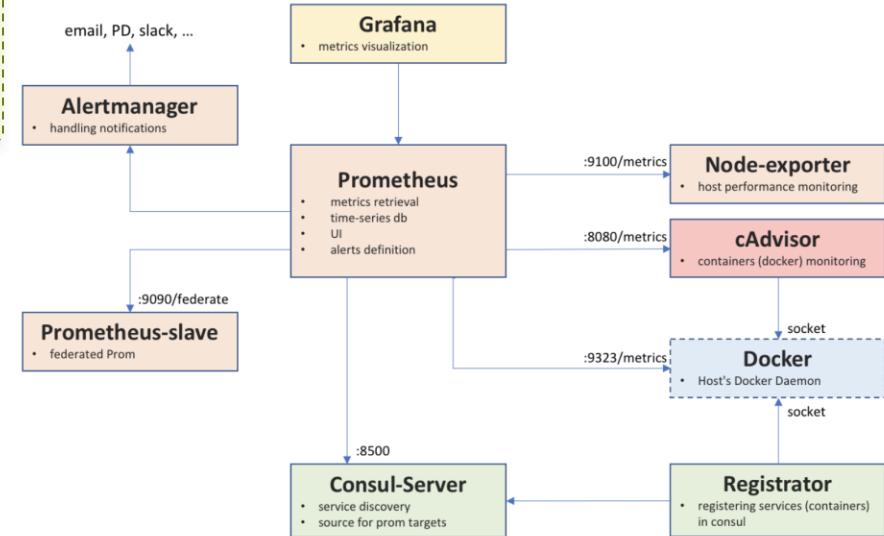
User	PID	PPID	Start Time	CPU %	MEM %	RSS	Virtual Size	Status	Running Time	Command	Container
root	9,694	9,677	18:59	5.00	13.70	67.09 MiB	443.49 MiB	Ssl	00:07:48	cadvisor	/docker/76c9637d77b036
root	1,003	1	10:09	1.10	9.40	46.07 MiB	1.46 GiB	Ssl	00:07:58	dockerd	/system.slice/docker
root	35	2	10:08	1.00	0.00	0.00 B	0.00 B	S	00:07:04	kswapd0	
root	1,267	1,003	10:09	0.30	2.50	12.64 MiB	1007.03 MiB	Ssl	00:02:23	docker-containe	/system.slice/docker
root	3	2	10:08	0.10	0.00	0.00 B	0.00 B	S	00:00:57	ksoftirqd/0	
root	9	2	10:08	0.10	0.00	0.00 B	0.00 B	S	00:01:13	rcu_sched	

# Collect Docker metrics with Prometheus

/etc/docker/daemon.json

```
{  
    "metrics-addr" : "0.0.0.0:9323",  
    "experimental" : true  
}
```

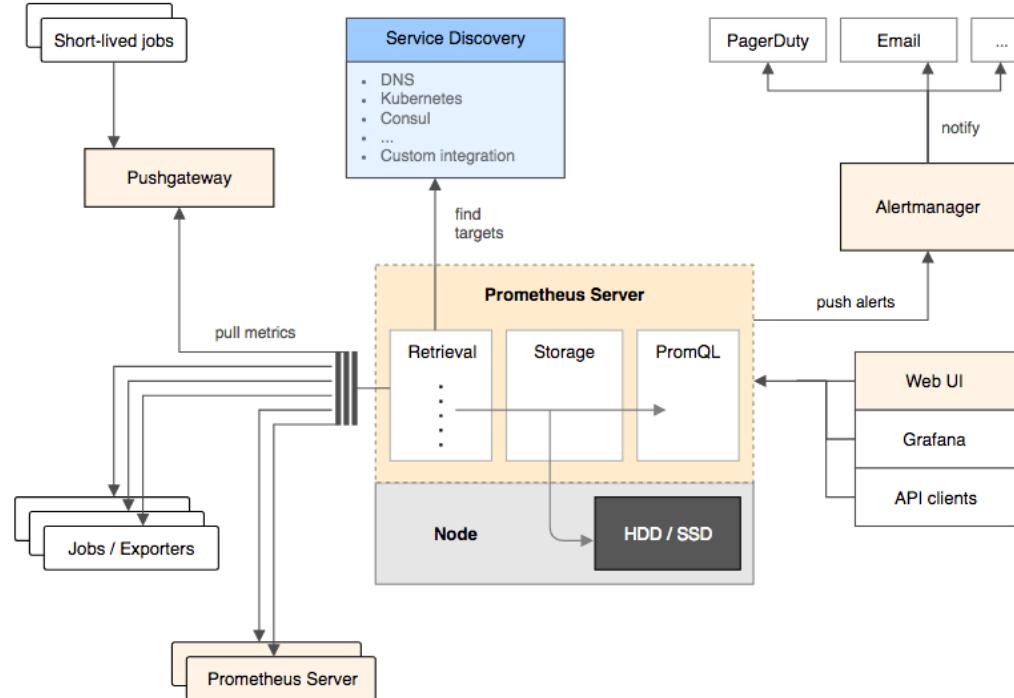
Detailed Example: [monitoring/prometheus](#)

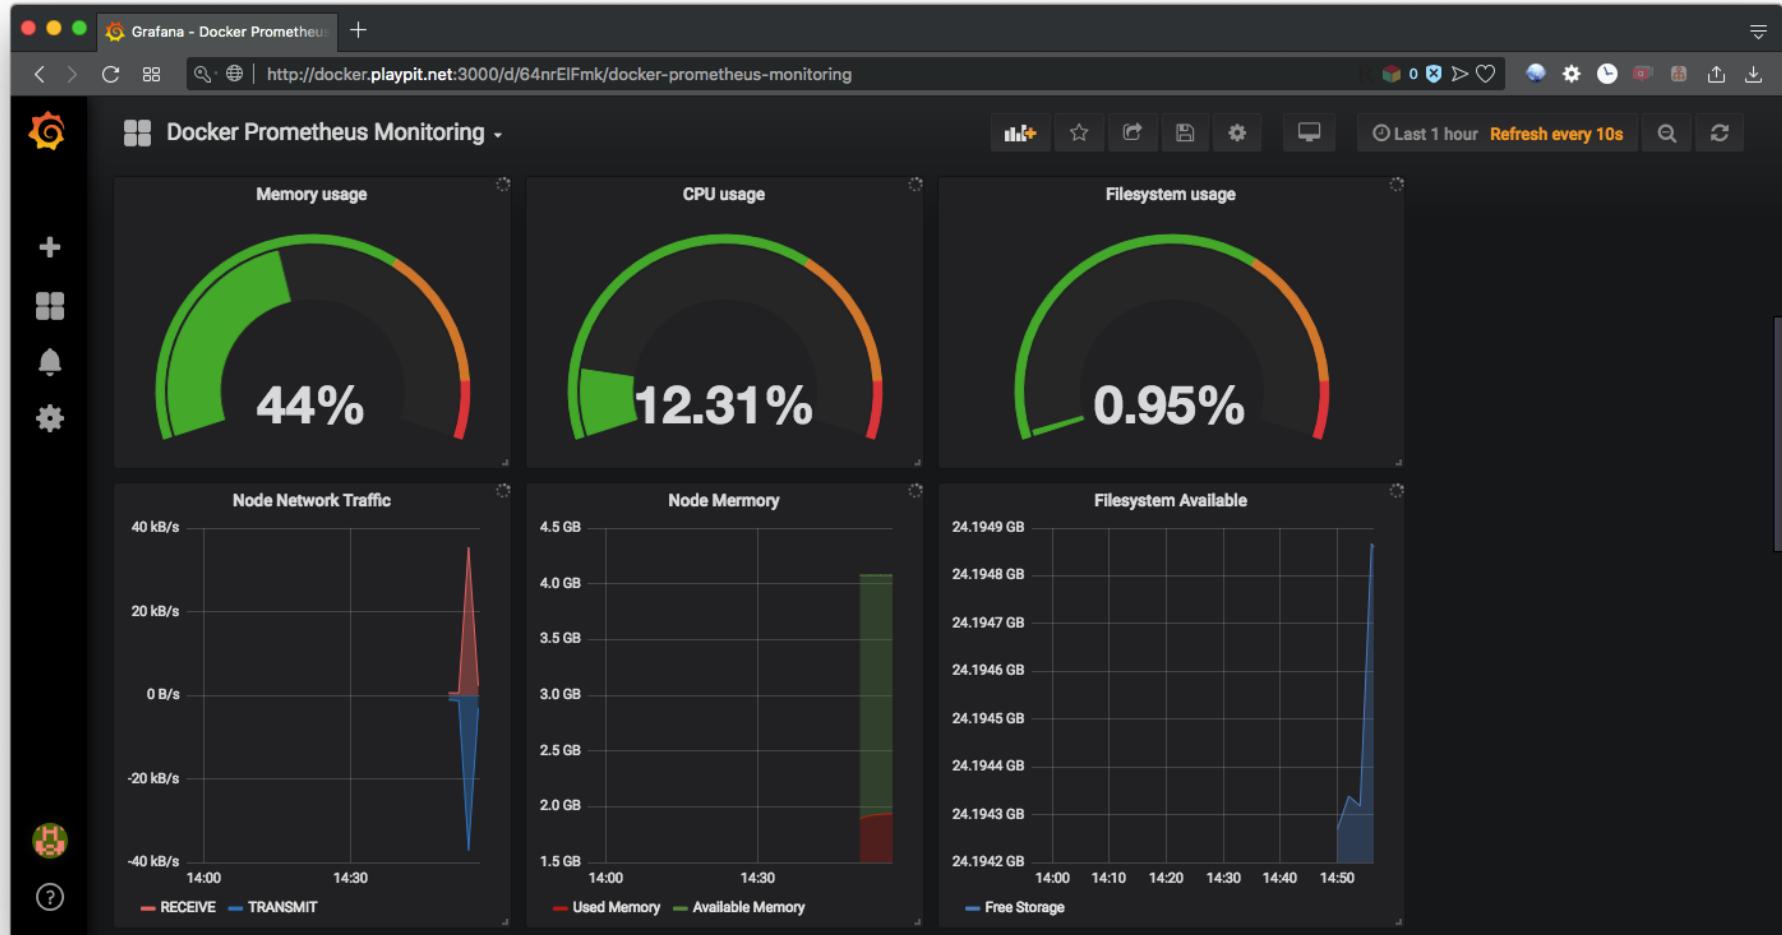


Prometheus Docker: <https://docs.docker.com/config/thirdparty/prometheus/>

Prometheus Project: <https://prometheus.io/>

# Prometheus Architecture





**Execute**

engine\_daemon\_network\_act ▾

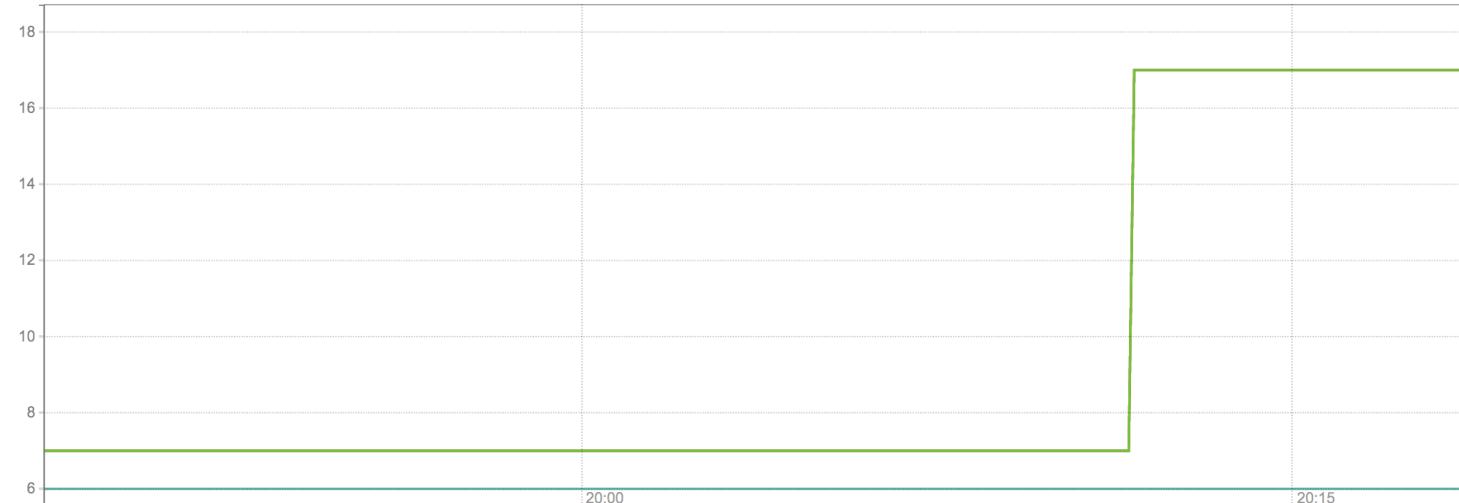
**Graph****Console**

- 30m +

◀ Until ▶

Res. (s)

□ stacked



- ✓ engine\_daemon\_network\_actions\_seconds\_count{action="release",instance="192.168.65.1:9323",job="docker"}
- ✓ engine\_daemon\_network\_actions\_seconds\_count{action="connect",instance="192.168.65.1:9323",job="docker"}
- ✓ engine\_daemon\_network\_actions\_seconds\_count{action="allocate",instance="192.168.65.1:9323",job="docker"}

# Useful Commands: Getting Info

Display system-wide information

```
# docker system info
```

Show docker disk usage

```
# docker system df
```

Show detailed information on space usage

```
# docker system df -v
```

Display total file sizes

```
# docker ps -s
```

Display resource usage statistics

```
# docker stats --no-stream
```

Show exited containers

```
# docker ps -f status=exited
```

Show unused images

```
# docker images -f dangling=true
```

Show unused volumes

```
# docker volume ls -f dangling=true
```

# Useful Commands: Cleaning Up

## # docker image prune

WARNING! This will remove all dangling images.  
Are you sure you want to continue? [y/N]  
you want to continue? [y/N]

## # docker container prune

WARNING! This will remove all stopped containers.  
Are you sure you want to continue? [y/N]

## # docker volume prune

WARNING! This will remove all volumes not used by at least one container.  
Are you sure you want to continue? [y/N]

## # docker network prune

WARNING! This will remove all networks not used by at least one container.  
Are you sure you want to continue? [y/N]

## # docker system prune

WARNING! This will remove:

- all stopped containers
- all networks not used by at least one container
- all dangling images
- all build cache

Are you sure you want to continue? [y/N]



## Containers Vulnerabilities

- Checking Containers

# Docker Security

---

## Docker vulnerabilities and threats to battle

- [Docker host and kernel security](#)
- [Docker container breakout](#)
- [Container image authenticity](#)
- [Container resource abuse](#)
- [Docker security vulnerabilities present in the static image](#)
- [Docker credentials and secrets](#)
- [Docker runtime security monitoring](#)

# Docker Image Vulnerabilities

[https://hub.docker.com/\\_/centos?tab=tags](https://hub.docker.com/_/centos?tab=tags)

Tags (33)



centos

By Docker

latest 75 MB

Last update: a month ago

This image has vulnerabilities



centos7.6.1810 75 MB

Last update: a month ago

This image has vulnerabilities



centos7.5.1804 75 MB

Last update: a month ago

This image has vulnerabilities



centos7.4.1708 73 MB

Last update: a month ago

This image has vulnerabilities



[https://hub.docker.com/\\_/ubuntu?tab=tags](https://hub.docker.com/_/ubuntu?tab=tags)

Tags (319)



ubuntu



Docker Official Images

Ubuntu is a Debian-based Linux operating system based on free software.

xenial-20190515 44 MB

Last update: a month ago

This image has vulnerabilities



xenial 44 MB

Last update: a month ago

This image has vulnerabilities



trusty-20190515 67 MB

Last update: a month ago

This image has vulnerabilities



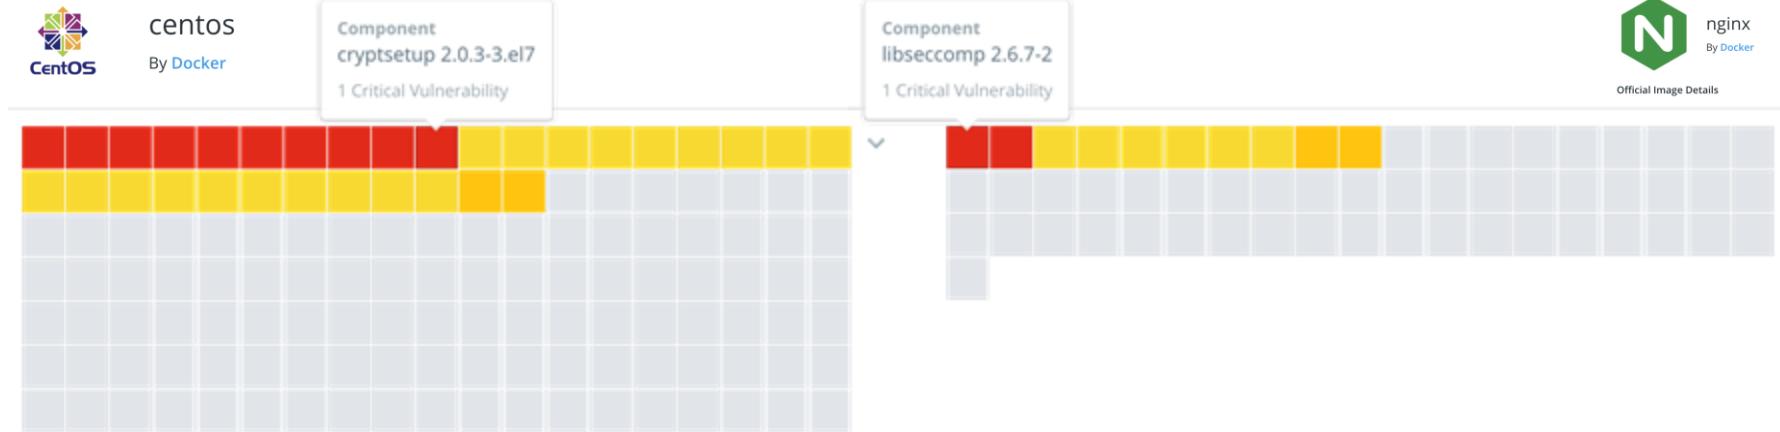
trusty 67 MB

Last update: a month ago

This image has vulnerabilities



# Docker Image Vulnerabilities



# Container Image Scanning Open-source Tools

---

There are several Docker image scanning tools available, and some of the most popular include:

- **Anchore Engine**: Anchore Engine is an open source image scanning tool. Provides a centralized service for inspection, analysis and applies user-defined acceptance policies to allow automated validation and certification of container images.
- **CoreOS/Clair**: An open source project for the static analysis of vulnerabilities in application containers (currently including appc/Rkt and Docker).
- **Vuls.io**: Agent-less Linux vulnerability scanner based on information from NVD, OVAL, etc. It has some container image support, although is not a container specific tool.
- **OpenScap**: Suite of automated audit tools to examine the configuration and known vulnerabilities in your software, following the NIST-certified Security Content Automation Protocol (SCAP). Not container specific again, but does include some level of support.
- etc

<https://sysdig.com/blog/20-docker-security-tools/>

# Anchore Engine

<https://github.com/anchore/anchore-engine>

<https://github.com/anchore/anchore-engine#configuration>

<https://github.com/anchore/anchore-engine#running-anchore-engine-using-docker-compose>

<https://sysdig.com/blog/container-security-docker-image-scanning/>

```
● ● ● sbeliakou@HOST:~/aevolume (ssh)
sbeliakou:~/aevolume $ docker-compose ps
          Name           Command       State        Ports
-----+-----+-----+-----+
aevolume_anchore-db_1      docker-entrypoint.sh postgres    Up            5432/tcp
aevolume_engine-analyzer_1  docker-entrypoint.sh anch ...  Up (healthy)  8228/tcp
aevolume_engine-api_1       docker-entrypoint.sh anch ...  Up (healthy)  0.0.0.0:8228->8228/tcp
aevolume_engine-catalog_1   docker-entrypoint.sh anch ...  Up (healthy)  8228/tcp
aevolume_engine-policy-engine_1  docker-entrypoint.sh anch ...  Up (healthy)  8228/tcp
aevolume_engine-simpleq_1   docker-entrypoint.sh anch ...  Up (healthy)  8228/tcp
sbeliakou:~/aevolume $
sbeliakou:~/aevolume $ anchore-cli system status
Service simplequeue (anchore-quickstart, http://engine-simpleq:8228): up
Service policy_engine (anchore-quickstart, http://engine-policy-engine:8228): up
Service apixext (anchore-quickstart, http://engine-api:8228): up
Service catalog (anchore-quickstart, http://engine-catalog:8228): up
Service analyzer (anchore-quickstart, http://engine-analyzer:8228): up

Engine DB Version: 0.0.10
Engine Code Version: 0.4.0

sbeliakou:~/aevolume $ █
```

```
$ yum install epel-release  
$ yum install python-pip  
$ pip install anchorecli
```

```
2. sbeliakou@HOST:~/aevolume (ssh)  
sbeliakou:~/aevolume $ anchore-cli image add docker.io/library/ubuntu:19.04  
Image Digest: sha256:887a7c706e88e4569494c9f470a53217a2705c9ddb3f8f75dc9062ece9e0bd36  
Parent Digest: sha256:1f070d23781c4e0aa3841ae9d7494885283e199e61eddec8337cd1895f864d67  
Analysis Status: not_analyzed  
Image Type: docker  
Image ID: 9b17fc7d68486409b6adf12e02fbf19b509714e5e7e2993e38d05ef47a29bfcc  
Dockerfile Mode: None  
Distro: None  
Distro Version: None  
Size: None  
Architecture: None  
Layer Count: None  
  
Full Tag: docker.io/library/ubuntu:19.04  
  
sbeliakou:~/aevolume $ anchore-cli image list  
Full Tag Image Digest Analysis Status  
docker.io/centos:7 sha256:ca58fe458b8d94bc6e3072f1cfbd334855858e05e1fd633aa07cf7f82b048e66 not_analyzed  
docker.io/jenkins/jenkins:latest sha256:bb0dce6f9887619e01bbeff11b6f99a9470138f65ab570f7813d25029f43cb1e analyzed  
docker.io/jenkins/jenkins:lts sha256:08bdd27b066e4ec6bfa1228876551e939df1aeaa9d878e6ddb2b183cd208dc2b analyzing  
docker.io/library/python:2.7 sha256:3079b0d54fd139626b117f01ce61fb3e84e37b73c5334d80b47cdd196a0e5036 analyzed  
docker.io/library/python:3.5 sha256:8e4a6308dbdf20a6c46b79449d1eb7c191186c790233fdf247e5f40aea14438c not_analyzed  
docker.io/library/ubuntu:19.04 sha256:887a7c706e88e4569494c9f470a53217a2705c9ddb3f8f75dc9062ece9e0bd36 not_analyzed  
sbeliakou:~/aevolume $
```

# Anchore Vulnerabilities Feeds

```
sbeliakou:~/aevolume $ anchore-cli system feeds list
2. sbeliakou@HOST:~/aevolume (ssh)
Feed          Group           LastSync          RecordCount
vulnerabilities alpine:3.3  2019-06-10T19:12:10.366919 457
vulnerabilities alpine:3.4  2019-06-10T19:12:10.911830 681
vulnerabilities alpine:3.5  2019-06-10T19:12:11.480844 875
vulnerabilities alpine:3.6  2019-06-10T19:12:12.113908 1051
vulnerabilities alpine:3.7  2019-06-10T19:12:12.532327 1125
vulnerabilities alpine:3.8  2019-06-10T19:12:12.927154 1220
vulnerabilities alpine:3.9  2019-06-10T19:12:13.302864 1218
vulnerabilities amzn:2     2019-06-10T19:12:13.837369 178
vulnerabilities centos:5   2019-06-10T19:12:14.376749 1323
vulnerabilities centos:6   2019-06-10T19:12:15.015396 1333
vulnerabilities centos:7   2019-06-10T19:12:15.429188 793
vulnerabilities debian:10   2019-06-10T18:27:51.278648 20344
vulnerabilities debian:7    2019-06-10T18:31:48.254963 20455
vulnerabilities debian:8    2019-06-10T18:36:28.647939 21767
vulnerabilities debian:9    2019-06-10T18:41:23.098457 20556
vulnerabilities debian:unstable 2019-06-10T18:46:11.455132 21236
vulnerabilities ol:5      2019-06-10T18:46:53.848234 1233
vulnerabilities ol:6      2019-06-10T18:47:47.408662 1417
vulnerabilities ol:7      2019-06-10T18:48:31.681759 915
vulnerabilities ubuntu:12.04 2019-06-10T18:51:35.328688 14948
vulnerabilities ubuntu:12.10 2019-06-10T18:52:40.950951 5652
vulnerabilities ubuntu:13.04 2019-06-10T18:53:25.228455 4127
vulnerabilities ubuntu:14.04 2019-06-10T18:57:07.312641 18679
vulnerabilities ubuntu:14.10 2019-06-10T18:57:59.884495 4456
vulnerabilities ubuntu:15.04 2019-06-10T18:59:05.891355 5789
vulnerabilities ubuntu:15.10 2019-06-10T19:00:27.227882 6513
vulnerabilities ubuntu:16.04 2019-06-10T19:03:37.687545 15780
vulnerabilities ubuntu:16.10 2019-06-10T19:05:08.256536 8647
vulnerabilities ubuntu:17.04 2019-06-10T19:06:49.328551 9157
vulnerabilities ubuntu:17.10 2019-06-10T19:08:09.246053 7935
vulnerabilities ubuntu:18.04 2019-06-10T19:09:48.429977 10032
vulnerabilities ubuntu:18.10 2019-06-10T19:11:06.915289 8119
vulnerabilities ubuntu:19.04 2019-06-10T19:12:09.953814 6571
sbeliakou:~/aevolume $
```

2. sbeliakou@HOST:~/aevolume (ssh)

Vulnerability ID	Package	Severity	Fix	Vulnerability URL
CVE-2017-14062	libidn11-1.33-1	High	None	<a href="https://security-tracker.debian.org/tracker/CVE-2017-14062">https://security-tracker.debian.org/tracker/CVE-2017-14062</a>
CVE-2017-17458	mercurial-4.0-1+deb9u1	High	None	<a href="https://security-tracker.debian.org/tracker/CVE-2017-17458">https://security-tracker.debian.org/tracker/CVE-2017-17458</a>
CVE-2017-17458	mercurial-common-4.0-1+deb9u1	High	None	<a href="https://security-tracker.debian.org/tracker/CVE-2017-17458">https://security-tracker.debian.org/tracker/CVE-2017-17458</a>
CVE-2018-13347	mercurial-4.0-1+deb9u1	High	None	<a href="https://security-tracker.debian.org/tracker/CVE-2018-13347">https://security-tracker.debian.org/tracker/CVE-2018-13347</a>
CVE-2018-13347	mercurial-common-4.0-1+deb9u1	High	None	<a href="https://security-tracker.debian.org/tracker/CVE-2018-13347">https://security-tracker.debian.org/tracker/CVE-2018-13347</a>
CVE-2019-8457	libssqlite3-0.3.16.2-5+deb9u1	High	None	<a href="https://security-tracker.debian.org/tracker/CVE-2019-8457">https://security-tracker.debian.org/tracker/CVE-2019-8457</a>
CVE-2018-16868	libgnutls30-3.5.8-5+deb9u4	Low	None	<a href="https://security-tracker.debian.org/tracker/CVE-2018-16868">https://security-tracker.debian.org/tracker/CVE-2018-16868</a>
CVE-2011-3389	libgnutls30-3.5.8-5+deb9u4	Medium	None	<a href="https://security-tracker.debian.org/tracker/CVE-2011-3389">https://security-tracker.debian.org/tracker/CVE-2011-3389</a>
CVE-2018-1000132	mercurial-4.0-1+deb9u1	Medium	None	<a href="https://security-tracker.debian.org/tracker/CVE-2018-1000132">https://security-tracker.debian.org/tracker/CVE-2018-1000132</a>
CVE-2018-1000132	mercurial-common-4.0-1+deb9u1	Medium	None	<a href="https://security-tracker.debian.org/tracker/CVE-2018-1000132">https://security-tracker.debian.org/tracker/CVE-2018-1000132</a>
CVE-2018-12404	libnss3-2:3.26.2-1.1+deb9u1	Medium	None	<a href="https://security-tracker.debian.org/tracker/CVE-2018-12404">https://security-tracker.debian.org/tracker/CVE-2018-12404</a>
CVE-2018-13346	mercurial-4.0-1+deb9u1	Medium	None	<a href="https://security-tracker.debian.org/tracker/CVE-2018-13346">https://security-tracker.debian.org/tracker/CVE-2018-13346</a>

2. sbeliakou@HOST:~/aevolume (ssh)

Vulnerability ID	Package	Severity	Fix	Vulnerability URL
RHSA-2019:0679	libssh2-1.4.3-12.el7	High	0:1.4.3-12.el7_6.2	<a href="https://access.redhat.com/errata/RHSA-2019:0679">https://access.redhat.com/errata/RHSA-2019:0679</a>
RHSA-2019:0710	python-2.7.5-76.el7	High	0:2.7.5-77.el7_6	<a href="https://access.redhat.com/errata/RHSA-2019:0710">https://access.redhat.com/errata/RHSA-2019:0710</a>
RHSA-2019:0710	python-libs-2.7.5-76.el7	High	0:2.7.5-77.el7_6	<a href="https://access.redhat.com/errata/RHSA-2019:0710">https://access.redhat.com/errata/RHSA-2019:0710</a>
RHSA-2019:1294	bind-license-9.9.4-73.el7_6	High	32:9.9.4-74.el7_6.1	<a href="https://access.redhat.com/errata/RHSA-2019:1294">https://access.redhat.com/errata/RHSA-2019:1294</a>
RHSA-2019:0483	openssl-libs-1.0.2k-16.el7	Medium	1:1.0.2k-16.el7_6.1	<a href="https://access.redhat.com/errata/RHSA-2019:0483">https://access.redhat.com/errata/RHSA-2019:0483</a>

sbeliakou:~/aevolume \$

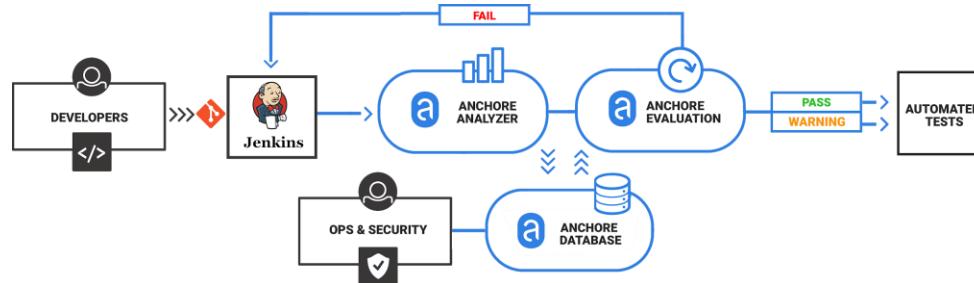
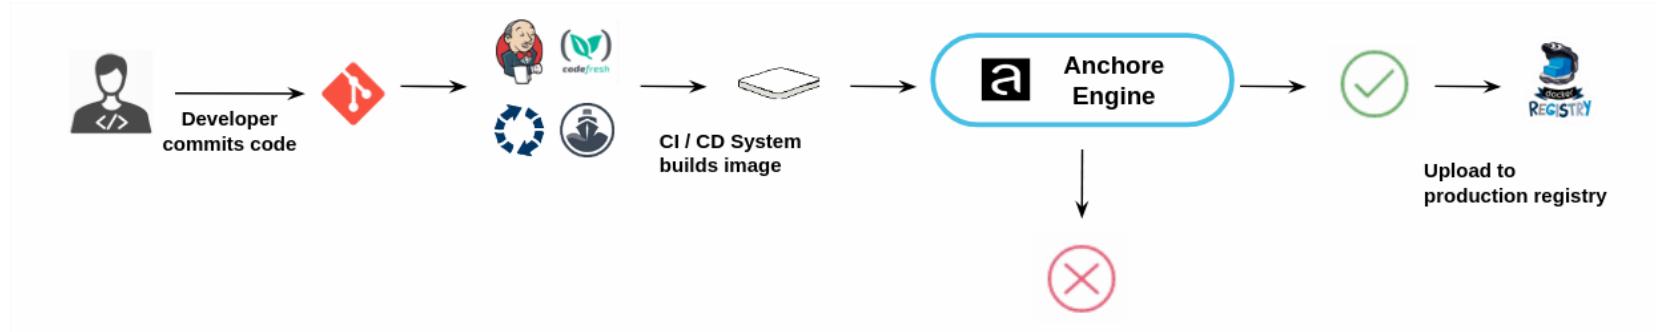
2. sbeliakou@HOST:~/aevolume (ssh)

Vulnerability ID	Package	Severity	Fix	Vulnerability URL
CVE-2013-4235	login-1:4.5-ubuntu2	Low	None	<a href="http://people.ubuntu.com/~ubuntu-security/cve/CVE-2013-4235">http://people.ubuntu.com/~ubuntu-security/cve/CVE-2013-4235</a>
CVE-2013-4235	passwd-1:4.5-ubuntu2	Low	None	<a href="http://people.ubuntu.com/~ubuntu-security/cve/CVE-2013-4235">http://people.ubuntu.com/~ubuntu-security/cve/CVE-2013-4235</a>
CVE-2015-8985	libc-bin-2.27-3ubuntu1	Low	None	<a href="http://people.ubuntu.com/~ubuntu-security/cve/CVE-2015-8985">http://people.ubuntu.com/~ubuntu-security/cve/CVE-2015-8985</a>
CVE-2015-8985	libc6-2.27-3ubuntu1	Low	None	<a href="http://people.ubuntu.com/~ubuntu-security/cve/CVE-2015-8985">http://people.ubuntu.com/~ubuntu-security/cve/CVE-2015-8985</a>
CVE-2016-2781	coreutils-8.28-1ubuntu1	Low	None	<a href="http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-2781">http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-2781</a>
CVE-2018-16868	libgnutls30-3.5.18-1ubuntu1	Low	None	<a href="http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-16868">http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-16868</a>
CVE-2018-16869	libhogweed4-3.4-1	Low	None	<a href="http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-16869">http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-16869</a>
CVE-2018-16869	libnettle6-3.4-1	Low	None	<a href="http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-16869">http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-16869</a>

2. sbeliakou@HOST:~/aevolume (ssh)

sbeliakou:~/aevolume \$ anchore-cli image vuln docker.io/library/python:2.7 all | wc -l  
1832  
sbeliakou:~/aevolume \$

# Docker Scanner with Jenkins



<https://wiki.jenkins.io/display/JENKINS/Anchore+Container+Image+Scanner+Plugin>  
<https://medium.com/accelero/container-security-with-anchore-8785efa644a5>

# Docker Scanner with Jenkins

Policy

## Anchore Policy Evaluation Summary

Show 10 entries

Search:

Repo Tag	Stop Actions	Warn Actions	Go Actions	Final Action
docker.io/library/ubuntu:latest	0	14	0	WARN

Showing 1 to 1 of 1 entries

Previous 1 Next

## Anchore Policy Evaluation Report

Show 10 entries

Search:

Image Id	Repo Tag	Trigger Id	Gate	Trigger	Check Output	Gate Action	Whitelisted
f975c50357489439eb9145dbfa16bb7cd06c02c31aa4df45c77de4d2baa4e232	docker.io/library/ubuntu:latest	b38090bac771995c5af3fc8c033b7d3d	dockerfilecheck	nohealthcheck	Dockerfile does not contain any HEALTHCHECK instructions	WARN	false
f975c50357489439eb9145dbfa16bb7cd06c02c31aa4df45c77de4d2baa4e232	docker.io/library/ubuntu:latest	CVE-2018-6829+gnupg	anchoresec	vulnmedium	MEDIUM Vulnerability found in package - gnupg (CVE-2018-6829 - http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-6829)	WARN	false
f975c50357489439eb9145dbfa16bb7cd06c02c31aa4df45c77de4d2baa4e232	docker.io/library/ubuntu:latest	CVE-2018-6829+gpgv	anchoresec	vulnmedium	MEDIUM Vulnerability found in package - gpgv (CVE-2018-6829 - http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-6829)	WARN	false

<https://wiki.jenkins.io/display/JENKINS/Anchore+Container+Image+Scanner+Plugin>

# Anchore CI Demos

<https://github.com/Btdohunter/ci-demos>



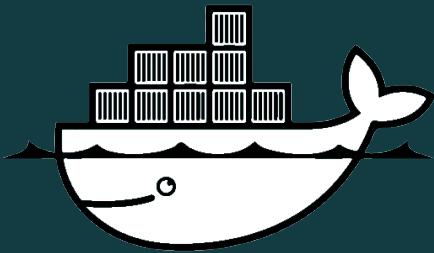
Jenkins

```
1 pipeline{
2     agent {
3         docker {
4             image 'docker:stable'
5         }
6     }
7     environment {
8         IMAGE_NAME = 'btdohunter/anchore-ci-demo'
9         IMAGE_TAG = 'jenkins'
10    }
11   stages {
12       stage('Build Image') {
13           steps {
14               sh 'docker build -t ${IMAGE_NAME}:ci .'
15           }
16       }
17       stage('Scan') {
18           steps {
19               sh 'apk add bash curl'
20               sh 'curl -s https://ci-tools.anchore.io/inline_scan-v0.3.3 | bash -s -- -d'
21           }
22       }
23       stage('Push Image') {
24           steps {
25               withDockerRegistry([credentialsId: "dockerhub-creds", url: ""]){
26                   sh 'docker tag ${IMAGE_NAME}:ci ${IMAGE_NAME}:${IMAGE_TAG}'
27                   sh 'docker push ${IMAGE_NAME}:${IMAGE_TAG}'
28               }
29           }
30       }
31   }
32 }
```



GitLab

```
1 variables:
2     IMAGE_NAME: ${CI_REGISTRY_IMAGE}/build:${CI_COMMIT_REF_SLUG}-${CI_COMMIT_SHA}
3
4 stages:
5     - build
6     - scan
7     - publish
8
9 container_build:
10    stage: build
11    image: docker:stable
12    services:
13        - docker:stable-dind
14
15 variables:
16     DOCKER_DRIVER: overlay2
17
18 script:
19     - echo "$CI_JOB_TOKEN" | docker login -u gitlab-ci-token --password-stdin "${CI_REGISTRY}"
20     - docker build -t "$IMAGE_NAME" .
21     - apk add bash curl
22     - curl -s https://ci-tools.anchore.io/inline_scan-v0.3.3 | bash -s ---r -t 500 "$IMAGE_NAME"
23     - docker push "$IMAGE_NAME"
24     - |
25         echo "Parsing anchore reports."
26         apk add jq
27         bash <<EOF
28         for f in anchore-reports/*; do
29             if [[ "$f" == "content-os" ]]; then
30                 printf "\n%s\n" "The following OS packages are installed on ${IMAGE_NAME}:"
31                 jq '[.content | sort_by(.package) | .[] | {package: .package, version: .version}]' $f || true
32             fi
33             if [[ "$f" == "vuln" ]]; then
34                 printf "\n%s\n" "The following vulnerabilities were found on ${IMAGE_NAME}:"
35                 jq '[.vulnerabilities | group_by(.package) | .[] | {package: .@[], vuln: [.].vuln}]]' $f || true
36             fi
37         done
38 EOF
```



**That's it for this training!**  
**Thank you for your attention!**

*Siarhei Beliakou,  
2019*